`bispectrum`: Selective G-Bispectra Made Practical

Anonymous Author(s)
Affiliation
Address
email

Abstract

Many machine-learning tasks are invariant under group actions: image classification under planar rotation $\mathrm{SO}(2)$, signal classification under translation, and spherical-image classification under 3D rotations $\mathrm{SO}(3)$. The G-bispectrum is a principled *complete* invariant (it preserves all signal information up to the group action) but its deployment has been hampered by quadratic cost and a patchwork of group-specific implementations. We present `bispectrum`, an open-source, fully unit-tested PyTorch library that computes *selective* G-bispectral invariants for seven different group/domain pairs as drop-in differentiable `nn.Modules`. Selectivity reduces the invariant size from $O(|G|^2)$ for the full bispectrum to $O(|G|)$ while remaining complete on generic signals. On the theoretical side, we contribute a selective $\mathrm{SO}(3)$ invariant on the sphere S^2 that achieves $\Theta(L^2)$ coefficients at band-limit L, matching the information-theoretic orbit-space lower bound $(L+1)^2 - 3$ instead of the $O(L^3)$ of the full bispectrum. We give a Kakarala-style conditional completeness theorem reducing the question to two cleanly named hypotheses (finite-band seed injectivity and uniform bootstrap-determinant nonvanishing), both supported by per-degree numerical certificates up to $\ell = 100$. We benchmark the library on three datasets across three group/domain pairs — PCam under C_8, OrganMNIST3D under the octahedral group, and Spherical MNIST under $\mathrm{SO}(3)$ — comparing against norm pooling, gated pooling, Fourier-ELU pooling, max pooling, and (non-equivariant) augmented CNN baselines. The library delivers near-exact rotation invariance (test-set rotation-accuracy std $\sigma_{\mathrm{rot}} < 10^{-3}$) and a data-efficiency advantage over incomplete alternatives in the low-data, moderate-capacity regime; the selective forward pass runs in sub-millisecond time on GPU for finite groups. Code is available at `https://anonymous.4open.science/r/bispectrum/`.

1 Introduction

Many machine-learning tasks are invariant under the action of a symmetry group G: histopathology classification under planar rotations C_n, signal classification under translations, volumetric organ classification under cube rotations O, and spherical-image classification under 3D rotations $\mathrm{SO}(3)$. Equivariant neural networks [Cohen and Welling, 2016, Weiler and Cesa, 2019] enforce these symmetries by construction in their feature maps, but at the point where a prediction must be made any task that does not depend on the group action — classification, regression, detection — still has to map the equivariant features to an *invariant* representation.

The *G-bispectrum* is a complex-valued vector of triple-correlation coefficients of the group Fourier transform. It has three properties that matter for ML: (i) **completeness** — it preserves all signal information up to the group action (the map is one-to-one on group orbits); (ii) **inversion** — the signal can in principle be recovered from the invariant up to its orbit; and (iii) **selectivity** — a

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

small $O(|G|)$ index subset of the $O(|G|^2)$ full bispectrum still suffices for completeness on generic signals [Mataigne et al., 2024].

Existing invariant pooling layers — norm pooling, gate-based pooling, max/sum aggregation [Weiler and Cesa, 2019] — are simple but *incomplete*: they collapse equivariant features to scalars and discard information that distinguishes signals in different orbits. The bispectrum offers a mathematically principled alternative [Giannakis, 1989, Kakarala, 2012] and has been used as a competitive pooling layer in equivariant networks [Sanborn et al., 2023, Mataigne et al., 2024].

Despite this, bispectral pooling has not been widely adopted, for both theoretical and computational reasons. Existing bispectral constructions are scattered across theory papers, domain-specific derivations, and small implementations, with no general-purpose PyTorch library covering the discrete and continuous symmetry groups relevant to modern geometric deep learning. Even for groups where selectivity [Mataigne et al., 2024] reduces the bispectrum from $O(|G|^2)$ to $O(|G|)$ scalars (and the per-pass cost from $O(|G|^3)$ to $O(|G|^2)$), a user still needs a numerically stable implementation, a consistent API, an inversion routine, tests, and benchmarked runtime. On the theoretical side, the selective construction of Mataigne et al. [2024] applies to finite groups; the continuous case of $\mathrm{SO}(3)$ acting on the sphere S^2 has remained open.

This paper presents `bispectrum`, a PyTorch [Paszke et al., 2019] library for computing bispectral invariants and, where available, their inverses. The library supports the C_n, D_n, $\mathbb{T}^d$, disk, S^2, and octahedral domains as differentiable `nn.Modules` intended to plug into standard ML workflows. Beyond the software contribution, the paper makes one technical contribution that arose during implementation: a selective $\mathrm{SO}(3)$ bispectrum on S^2 with $\Theta(L^2)$ coefficients (Section 6).

Contributions.

1. **Software contribution.** A PyTorch library implementing selective G-bispectra for *seven different group/domain pairs* (a *domain* is the underlying space on which the group acts: C_n on C_n, $\mathrm{SO}(2)$ on S^1, $\mathbb{T}^d$, D_n, $\mathrm{SO}(2)$ on the disk, $\mathrm{SO}(3)$ on S^2, and the octahedral group O) as drop-in differentiable `nn.Modules` with a consistent API, GPU-optimized forward passes ($<$1 ms for finite groups), precomputed CG buffers, and a fully unit-tested codebase (96.8% line coverage).
2. **Theoretical contribution: selective $\mathrm{SO}(3)$ bispectrum on S^2.** We construct an invariant on S^2 that uses $\Theta(L^2)$ scalar coefficients (matching the orbit-space dimension $(L+1)^2 - 3$ up to a constant factor) instead of the $O(L^3)$ of the full bispectrum. We give a Kakarala-style *conditional* completeness theorem reducing the question to two named hypotheses — finite-band seed injectivity (H1, motivated by Kakarala [1992, 2012] and Bandeira et al. [2017]) and uniform bootstrap-determinant nonvanishing (H2) — both supported by per-degree numerical certificates at a deterministic witness for every $\ell \leq 100$; the verification script is uniform in L and extends to larger band-limits without changes to the structural argument (Section 6, Appendix A).
3. **Empirical validation.** A benchmark of G-bispectral pooling across three ML tasks — PCam under C_8, OrganMNIST3D under the octahedral group O, and Spherical MNIST under $\mathrm{SO}(3)$ — against norm pooling, gated pooling, Fourier-ELU pooling, max pooling, and (non-equivariant) augmented-CNN baselines. Complete bispectral invariants provide a data-efficiency advantage in the low-data, moderate-capacity regime while converging to incomplete alternatives at high capacity.

Together, these contributions provide the first general-purpose, drop-in differentiable library for G-bispectral invariants and the first practical extension of the selective construction from finite groups to $\mathrm{SO}(3)$ on S^2.

2 Mathematical Background

Notation. A *group* G is a set with an associative product, identity, and inverse (e.g. rotations of the plane, translations of an image grid). G is *abelian* if elements commute (e.g. cyclic groups C_n, the torus $\mathbb{T}^d$) and *non-abelian* otherwise (e.g. $\mathrm{SO}(3)$, the dihedral groups D_n, the octahedral group O). A *representation* ρ assigns to each $g \in G$ a unitary matrix $\rho(g)$ such that $\rho(gh) = \rho(g)\rho(h)$; an

Table 1: Invariant nonlinearities for equivariant CNNs. "Exact"* means the invariant has no discretisation/aliasing error; "Complete" means it separates orbits; the coefficient column counts the size of the output for a signal with $|G|$ Fourier samples. "Genericity assumption"† means completeness holds outside a Lebesgue-measure-zero algebraic subvariety of the signal space. The selective bispectrum is the only method that is simultaneously exact, complete, and linear-cost for finite groups.

Nonlinearity	Exact?*	Complete?	Coefficients	Forward cost	Key limitation
Norm	Yes	No	$O(\|G\|)$	$O(\|G\|)$	Discards phase
Gate	Yes	No	$O(\|G\|)$	$O(\|G\|)$	Scalar bottleneck
Fourier pointwise	No	Yes‡	$O(\|G\|)$	$O(\|G\| \log \|G\|)$	Aliasing
Full bispectrum	Yes	Yes	$O(\|G\|^2)$	$O(\|G\|^3)$	Quadratic cost
Selective bispectrum	**Yes**	**Yes**	$\mathbf{O(\|G\|)}$	$\mathbf{O(\|G\|^2)}$	Genericity assumption†

‡Complete pre-aliasing; the pointwise nonlinearity in the forward pass introduces aliasing that may break completeness in practice.

irreducible representation (*irrep*) admits no nontrivial invariant subspace. The *dual* $\hat{G}$ denotes the set of equivalence classes of irreps; for finite G, $|\hat{G}| = O(|G|)$. A *group action* of G on a domain X is a map $G \times X \to X$, $(g, x) \mapsto g \cdot x$, compatible with the group product; it induces an action on functions $f : X \to \mathbb{R}$ by $(f \circ g)(x) := f(g \cdot x)$. Two signals are in the same *orbit* if one is the image of the other under some $g \in G$; an *invariant* of f is a function constant on each orbit. For compact G there is a unique normalised translation-invariant measure $d\mu$ called the *Haar measure*; finite G is the special case where $d\mu$ is the counting measure. "*Lebesgue-almost-every* f" means "every f outside a measure-zero algebraic subvariety of the signal space." The groups and domains used in this paper are: the cyclic group C_n (discrete planar rotations), the dihedral group D_n (rotations and reflections of a regular n-gon), the torus $\mathbb{T}^d = (S^1)^d$ (multi-channel periodic signals), $\mathrm{SO}(2)$ acting on the disk (continuous planar rotations of a 2D image), $\mathrm{SO}(3)$ acting on the sphere S^2 (continuous 3D rotations of a spherical signal), and the chiral octahedral group O (24 proper rotations of the cube, acting on 3D voxel grids).

2.1 Equivariant features and invariant maps

A neural network layer Φ is *G-equivariant* if its output transforms predictably under the group action: $[\pi_{\mathrm{out}}(g) \circ \Phi](f) = \Phi(\pi_{\mathrm{in}}(g)f)$ for all $g \in G$ [Cohen and Welling, 2016, Weiler and Cesa, 2019]. Here f is an input signal (a function on the domain X), $\pi_{\mathrm{in}}, \pi_{\mathrm{out}}$ are the (representations of the) G-actions on the input and output feature spaces of Φ, and the composition $\pi_{\mathrm{out}}(g) \circ \Phi$ is the action of g on the output side. Stacking equivariant layers preserves equivariance, but any task that requires a prediction independent of the group action—classification, regression, detection—must eventually map equivariant features to *invariant* ones (constant on each orbit).

Table 1 summarises the trade-offs between exactness, completeness, and computational cost across existing invariant pooling layers. Norm pooling and gated nonlinearities are exact invariants but *incomplete*: they discard the phase of Fourier blocks, collapsing distinct orbits. Fourier-space pointwise nonlinearities preserve completeness but sacrifice exactness through aliasing [Weiler and Cesa, 2019]. The full bispectrum is exact and complete but costs $O(|G|^2)$ coefficients. The *selective* bispectrum (introduced below) preserves both exactness and completeness while using only $O(|G|)$ coefficients on generic signals.

2.2 The G-bispectrum and selective bispectra

Let G be a compact group with dual $\hat{G}$ and irreps $\{\rho_i\}_{i \in \hat{G}}$ of dimensions d_i. The *group Fourier transform* [Diaconis and Rockmore, 1990] of a signal $f : G \to \mathbb{C}$ is

$$\hat{f}(\rho_i) = \int_G f(g)\, \rho_i(g)\, d\mu(g) \quad \in \mathbb{C}^{d_i \times d_i}, \tag{1}$$

which specialises to $\sum_{g \in G} f(g)\, \rho_i(g)$ when G is finite (counting measure). The Clebsch–Gordan (CG) decomposition $\rho_i \otimes \rho_j \cong \bigoplus_{k \in \hat{G}} m_{ij}^k \rho_k$ expresses the tensor product of two irreps as a direct sum of irreps (with integer multiplicities $m_{ij}^k \geq 0$); the change-of-basis is realised by the *CG isom-*

etry $C_{ij}^k : V_{\rho_k} \to V_{\rho_i} \otimes V_{\rho_j}$, projecting the tensor product onto its ρ_k-isotypic component. For each admissible triple (i, j, k) with $m_{ij}^k > 0$, the *bispectrum* is the scalar

$$\beta_{i,j}^k(f) := \left\langle \, (C_{ij}^k)^\dagger \left(\hat{f}(\rho_i) \otimes \hat{f}(\rho_j)\right), \, \hat{f}(\rho_k) \right\rangle. \tag{2}$$

Each $\beta_{i,j}^k$ is invariant under the action $f \mapsto f \circ g$ and, taken over all admissible triples, the bispectrum $\{\beta_{i,j}^k\}_{(i,j,k)}$ is a *complete* invariant: it separates orbits for generic signals [Kakarala, 2012]. For finite G, $|\hat{G}| = O(|G|)$, so the full bispectrum has $O(|G|^2)$ scalar coefficients (and is naturally indexed by $(i, j) \in \hat{G}^2$ with k ranging over the CG-allowed irreps). Equation (2) is stated for G acting on itself; for homogeneous spaces of the form G/H (e.g. $\mathrm{SO}(3)$ on S^2, or $\mathrm{SO}(2)$ on the disk), the same construction applies after expanding f in a basis adapted to G/H (spherical harmonics for S^2, Bessel modes on the disk). The explicit specialisation for S^2 is in Definition 5.

Scalar vs. matrix (vector) bispectrum. Throughout this paper the term *scalar bispectrum* refers to the single complex number $\beta_{i,j}^k(f)$ obtained from the inner product in Eq. (2)—i.e., the CG-contracted scalar associated with each admissible triple. In the literature one also encounters the *matrix* or *vector* bispectrum, which keeps the un-contracted matrix $(C_{ij}^k)^\dagger(\hat{f}(\rho_i) \otimes \hat{f}(\rho_j))$ before pairing it with $\hat{f}(\rho_k)$. The two are related by a single trace/inner-product step; the scalar form is what the library outputs and what we analyse for $\mathrm{SO}(3)$ on S^2.

Selective bispectra. Mataigne et al. [2024] show that for finite groups, a breadth-first search (BFS) over the Kronecker product table of $\hat{G}$ selects $O(|G|)$ bispectral coefficients $\beta_{i,j}^k$ that remain complete; Myers and Miolane [2025] extend this to the disk. This reduction from quadratic to linear is the key to practical deployment: for example, the bispectrum of the octahedral group drops from 576 to 172 coefficients.

3 Related Work

Equivariant network libraries. `escnn` [Cesa et al., 2022] and `e3nn` [Geiger and Smidt, 2022] provide general-purpose equivariant convolution layers for 2D and 3D groups respectively, and `lie-nn` [Batatia et al., 2023] provides general Lie-group representation primitives (irreducible representations, Clebsch–Gordan decompositions, tensor products) on top of which custom equivariant layers can be built. These libraries implement equivariant *feature maps* but leave the choice of invariant pooling to the user (typically norm or gate). Our library is complementary: it provides the invariant map as a drop-in `nn.Module` that plugs into architectures built with any of these frameworks.

Bispectral approaches in neural networks. Sanborn et al. [2023] introduced bispectral neural networks (BNNs), using the full G-bispectrum as an invariant nonlinearity for C_n and D_n. Oreiller et al. [2022] applied bispectral CNNs to medical imaging with a fixed set of hand-selected triples. Both use the full (quadratic-cost) bispectrum and provide implementations tied to specific architectures. Our library generalizes to seven group/domain pairs, defaults to the $O(|G|)$ selective reduction [Mataigne et al., 2024], and exposes a framework-agnostic API. Chevalley et al. [2024] explore bispectral signatures for data characterisation, using bispectra as descriptors rather than learnable pooling primitives. Outside the equivariant-network setting, `PyBispectra` [Binns et al., 2025] computes bispectrum and bicoherence statistics of one-dimensional time-series signals for electrophysiology applications such as phase-amplitude coupling and time-delay estimation; the underlying 1D bispectrum is translation-invariant by construction, but the toolbox is not designed as a neural-network pooling layer.

Bispectra on homogeneous and non-homogeneous spaces. The selective G-bispectrum construction [Mataigne et al., 2024] reduces the cost of the G-bispectrum from $O(|G|^2)$ to $O(|G|)$ scalars on finite groups. Myers and Miolane [2025] extend the selective construction to the disk, where $\mathrm{SO}(2)$ acts on a non-homogeneous space. Kakarala [2012] treats the $\mathrm{SO}(3)$-on-$\mathrm{SO}(3)$ bispectrum and the closely related sphere case via Wigner-D irreps. The matrix (vector) bispectrum on S^2 is the homogeneous-space form on which our scalar contractions and CG power augmentation are built.

Table 2: Library profiling summary. Timings are median wall-clock on a single NVIDIA H100 80 GB GPU. "Coefs" = number of scalar coefficients (Sel. = selective, Full = full bispectrum); "Fwd" = forward-pass time in milliseconds; "Inv" = inversion time in milliseconds. "–" = not implemented. For continuous groups, $|G| = \infty$; coefficient counts refer to the band-limited representation at the tested resolution. *`SO3onS2` benchmarked at band-limit $\ell_{\max}$=16 (we use $\ell_{\max}$ and L interchangeably to denote the band-limit; see Section 6).

Module	G	$\|G\|$	Coefs		Fwd (ms)		Inv (ms)
			Sel.	Full	Sel.	Full	Sel.
`CnonCn`	C_{128}	128	128	8256	0.15	9.17	8.56
`TorusOnTorus`	$C_{32}\times C_{32}$	1024	1024	524800	0.08	0.71	68.02
`DnonDn`	D_{32}	64	245	–	2.52	–	42.33
`SO2onDisk`	$\mathrm{SO}(2)$	∞	411	–	1.94	–	17.69
`SO3onS2`*	$\mathrm{SO}(3)$	∞	430	–	26.54	–	–
`OctaonOcta`	O	24	172	–	0.70	–	–

CG product nonlinearities and tensor products. Tensor field networks [Thomas et al., 2018] and MACE [Batatia et al., 2022] use iterated Clebsch–Gordan products to build equivariant interactions. These products are exact and complete in principle, but each layer incurs $O(L^3)$ cost per CG contraction. Xie et al. [2025] analyse the *price of freedom* in tensor product architectures, showing that maintaining full expressivity through CG products becomes expensive at high angular momenta. The selective bispectrum achieves completeness (conditionally on H1 and H2 for $\mathrm{SO}(3)$ on S^2, see Section 6) in a single $O(L^2)$ pass for $\mathrm{SO}(3)$ or $O(|G|)$ pass for finite groups, avoiding iterated contractions.

Invariant theory and scattering. The wavelet scattering transform [Mallat, 2012] provides translation-invariant representations via iterated modulus operations; Esteves et al. [2018] extend this to the sphere. Unlike the G-bispectrum, scattering representations are stable but generally incomplete—they discard inter-scale phase information. The G-bispectrum's completeness guarantee [Giannakis, 1989, Kakarala, 2012] is its distinguishing theoretical advantage.

4 The `bispectrum` Library

4.1 Design and implementation

Every group/domain pair is exposed as a `torch.nn.Module` named `{Group}on{Domain}` (e.g. `SO3onS2`). Each module exposes a uniform API: `forward(f)` returns the selective bispectral invariants of f; `fourier(f)` returns the underlying group Fourier coefficients; and, where available, `invert(beta)` returns the reconstructed signal up to group-action indeterminacy. Modules accept raw real-valued inputs, handle the Fourier transform internally, and default to $O(|G|)$ selective coefficients (`selective=False` gives the full $O(|G|^2)$ set). Clebsch–Gordan (CG) matrices, discrete Fourier transform (DFT) kernels, and Bessel roots (used for the radial discretisation of the disk in `SO2onDisk`) are precomputed at construction and stored as non-learnable buffers; this lets `.to(device)` migrate the module transparently and ensures that forward passes incur zero recomputation. The only dependencies are PyTorch [Paszke et al., 2019], NumPy [Harris et al., 2020], and `torch_harmonics` [Bonev et al., 2023] (for `SO3onS2`). For $\mathrm{SO}(3)$ on S^2, reduced CG matrices are built in parallel and cached to disk; all forward methods are batch-native, using `torch.einsum`/`torch.bmm` for non-abelian modules and `torch.fft` for abelian ones. The test suite covers 96.8% of lines across all modules. A minimal usage example is in Appendix G.

4.2 Profiling

We profile the library on a single NVIDIA H100 80 GB GPU (PyTorch 2.10, batch = 16 for the forward pass and 4 for inversion, median wall-clock via `torch.utils.benchmark`); results are in Table 2 and Figure 1.

The selective forward pass is sub-millisecond for finite abelian modules (C_n: 0.15 ms, $\mathbb{T}^2$: 0.08 ms), under 2 ms for $\mathrm{SO}(2)$/disk (1.94 ms), and under 1 ms for the octahedral group (0.70 ms), making it

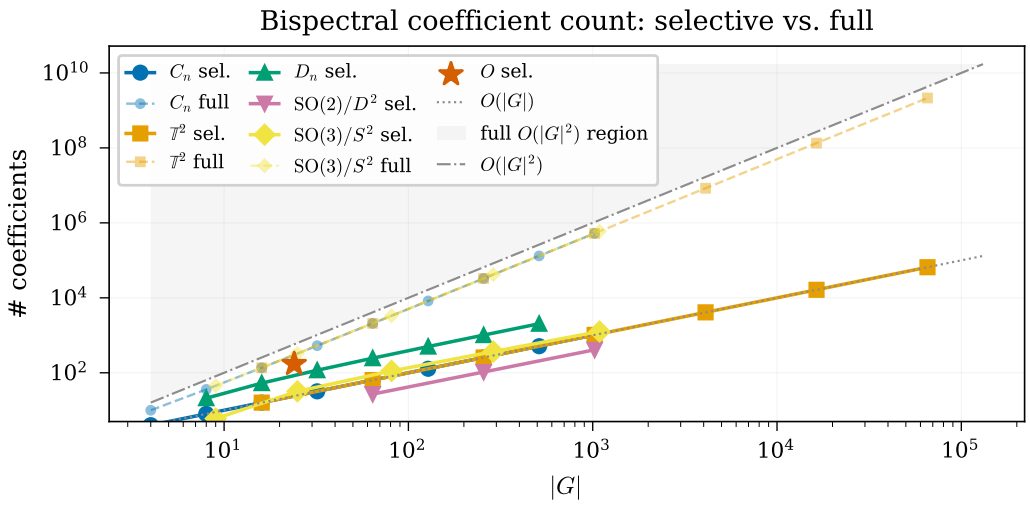

Figure 1: Bispectral coefficient count vs. group order $|G|$. Solid lines: selective ($O(|G|)$); dashed: full ($O(|G|^2)$). At $|G|$=1,024, selectivity yields a 512× reduction. Continuous groups (SO(2), SO(3)) are plotted at their band-limited signal dimension. The full-bispectrum curves for D_n and the disk are omitted because the corresponding backends are not yet implemented in the library.

negligible relative to backbone computation. SO(3) on S^2 at $\ell_{\max}$=16 takes 26.5 ms, dominated by the spherical harmonic transform (via `torch_harmonics` [Bonev et al., 2023]). Selective coefficient counts confirm the $O(|G|)$ scaling: C_{128} uses 128 coefficients vs. 8,256 full (64× reduction); $\mathbb{T}^2$ at $|G|$=1,024 uses 1,024 vs. 524,800 full (512×). Inversion is fast for abelian groups (9–68 ms). GPU throughput scales linearly with batch size, exceeding 10^7 samples/s for C_n and $\mathbb{T}^2$ (Appendix C).

5 Experiments

We validate the library on three tasks spanning discrete and continuous groups: 2D histopathology classification under C_8 symmetry (PCam, Section 5.1), 3D organ classification under octahedral symmetry (OrganMNIST3D, Section 5.2), and spherical digit classification under full SO(3) (Spherical MNIST, Section 5.3).

5.1 2D histopathology classification under cyclic symmetry

We evaluate five invariant pooling strategies on PatchCamelyon [Veeling et al., 2018] (PCam), 327K histopathology patches with binary metastasis labels. All equivariant models share a C_8-equivariant DenseNet [Huang et al., 2017] backbone following Veeling et al. [2018]; they differ only in the terminal invariant map: norm pooling, gated pooling, Fourier-ELU, and selective bispectrum via `CnonCn`(n=8). A non-equivariant standard CNN trained with data augmentation drawn uniformly from the symmetry group C_8 ("*augmented CNN*") serves as the non-equivariant baseline. We sweep growth rates to vary capacity (30K–1.6M params) and train on 10% of the PCam training set (an image-level subsample, not a group-orbit subsample, drawn before any group augmentation) with AdamW [Loshchilov and Hutter, 2019]. Full model descriptions are in Appendix D.

Results. Table 3 and Figure 2 present the results at matched ∼100K parameters.

At matched capacity (∼100K params, 10% data), all four equivariant pooling methods substantially outperform the standard augmented CNN (AUC 0.896). Fourier-ELU leads slightly (0.945), followed by NormReLU (0.942), bispectrum and gated (both 0.941). At 1% data, the bispectrum leads all methods (0.912), outperforming NormReLU (0.873) and Fourier-ELU (0.876). This data-efficiency advantage narrows as the training set grows, consistent with complete invariants providing a stronger inductive bias when data is scarce.

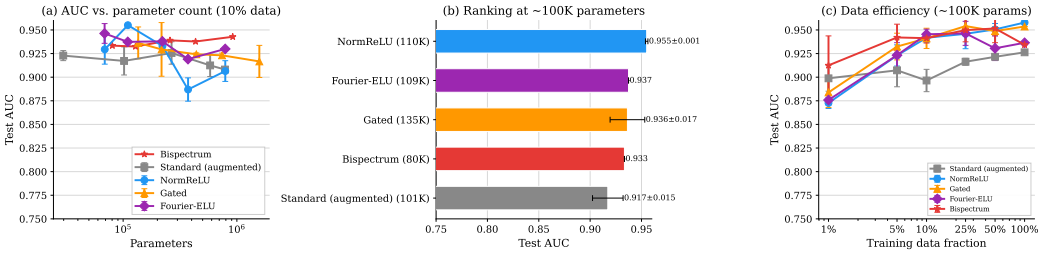

Figure 2: PCam Pareto analysis (`CnonCn` selective bispectrum, single seed, 10% training data). (a) Test AUC vs. parameter count across all five pooling variants. (b) Ranking at the ∼100K parameter budget. (c) Data efficiency at matched ∼100K parameters. Multi-seed statistics are in Table 3.

Table 3: PCam results at matched ∼100K parameters (10% training data, mean ± std over 3 seeds). All equivariant models share a C_8-equivariant DenseNet backbone.

Model	Params	Test AUC	Test ACC
Standard (augmented)	102K	.896 ± .010	.826 ± .012
Gated	136K	.941 ± .009	.870 ± .004
Selective Bispectrum (`CnonCn`)	128K	.941 ± .004	**.872 ± .006**
Fourier-ELU	110K	**.945 ± .004**	.855 ± .006
NormReLU	110K	.942 ± .004	.858 ± .007

5.2 3D organ classification under octahedral symmetry

To evaluate the `OctaonOcta` module for the chiral octahedral group O ($|O| = 24$), we design a controlled 3D experiment on OrganMNIST3D [Yang et al., 2023] (1,742 abdominal CT volumes, 28^3 voxels, 11 classes). All equivariant variants share an O-equivariant 3D ResNet backbone with exact voxel-permutation kernels. The four models compared are: a translation-equivariant standard 3D CNN trained with octahedral data augmentation (the "*augmented CNN*" baseline, 16K params); the same O-equivariant ResNet with terminal norm pool (375K params); with terminal max pool (374K params); and with terminal selective bispectrum via `OctaonOcta` (463K, 172 invariant scalars). Models are trained with AdamW [Loshchilov and Hutter, 2019] for 100 epochs with three seeds; rotation robustness is measured over all 24 octahedral rotations of the test set. Published baselines and wider-channel experiments are in Appendix E.

Results. Table 4 and Figure 3 report the key findings.

Rotation invariance. All equivariant models achieve $\sigma_{\text{rot}} \leq 0.0005$, while the augmented CNN drops ∼2.5 pp (percentage points) under rotation.

Completeness vs. lossy pooling. The bispectrum (72.6%) matches max pool (73.0%) and substantially outperforms norm pool (56.8%) and Regular SE(3) convolutions [Kuipers and Bekkers, 2023] (69.8%). At reduced data (Figure 3), the bispectrum's advantage widens: at 5% data it reaches 33.2% vs. 18.9% for max pool. At higher capacity (6M params, Appendix E), max pool overtakes the bispectrum, suggesting the completeness advantage is specific to the data-limited regime.

5.3 Spherical MNIST: SO(3) invariance on the sphere

The preceding experiments validate bispectral pooling for finite groups (C_8 and the octahedral group O). To evaluate the library's continuous group support, we test `SO3onS2`—the $\mathrm{SO}(3)$ bispectrum on the 2-sphere—on Spherical MNIST [Cohen et al., 2018], the standard benchmark for rotation invariance on S^2.

Setup. MNIST digits are projected onto the sphere via stereographic projection and sampled on a 64×128 equiangular grid. We compute the selective bispectrum at $\ell_{\max} = 15$, producing 384 com-

Table 4: OrganMNIST3D results (mean $\pm$ std over 3 seeds). Rot ACC and σ_{rot}: mean and std of accuracy across 24 octahedral test-set rotations.

Model	Params	Test ACC	Test AUC	Rot ACC	σ_{rot}
Standard 3D CNN	16K	$.601 \pm .017$	$.951 \pm .002$	.598	.0115
O-Equiv + Norm Pool	375K	$.568 \pm .099$	$.940 \pm .021$	.568	.0000
O-Equiv + Max Pool	374K	$.730 \pm .033$	$.972 \pm .005$	.730	.0000
O-Equiv + `OctaonOcta` Bispectrum	463K	$.726 \pm .027$	$.972 \pm .004$	.726	.0005†

†Nonzero σ_{rot} due to floating-point precision in CG contractions, not a theoretical invariance failure.

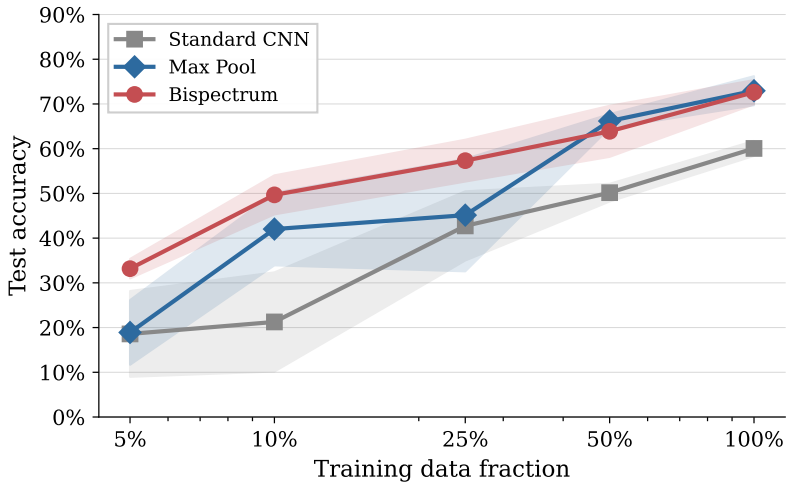

Figure 3: Test accuracy vs. training data fraction on OrganMNIST3D (`OctaonOcta` bispectrum, mean over 3 seeds; shaded bands: ± 1 std). The complete invariant map provides a consistent advantage in the low-data regime.

plex invariant coefficients (768 real features after taking log-magnitude), which are fed to a 3-layer MLP (232K parameters total). The conditional completeness theorem (Appendix A) is supported by per-degree numerical certificates for every $\ell \leq 100$, so $\ell_{\max} = 15$ is well within the verified range. Three baselines provide controls:

- **Power Spectrum + MLP** (11K params): computes only the rotation-invariant power spectrum $\|f_\ell\|^2$ per degree ℓ, discarding all phase coupling between harmonics. Invariant but *incomplete*.
- **Power Spectrum + MLP (matched)** (166K params): same incomplete invariant features but with a larger MLP (166K), isolating the effect of feature completeness from model capacity.
- **Standard CNN** (185K params): a conventional CNN operating directly on the equirectangular image. Not rotation-invariant.

Following Cohen et al. [2018], we evaluate under three protocols: train on non-rotated / test on non-rotated (NR/NR), train on rotated / test on rotated (R/R), and the most challenging train on non-rotated / test on rotated (NR/R). Each configuration is run with three seeds. Rotation robustness is assessed by applying 10 random $\mathrm{SO}(3)$ rotations to the entire test set and measuring accuracy variance.

Results. Table 5 and Figure 4 report the results.

Numerically exact $\mathrm{SO}(3)$ invariance. The bispectrum achieves near-identical accuracy across all three protocols ($\sigma_{\text{rot}} \sim 10^{-4}$), confirming true $\mathrm{SO}(3)$ invariance. The standard CNN collapses from 99.1% (NR/NR) to 13.8% (NR/R), showing that augmentation alone does not confer invariance.

Table 5: Spherical MNIST results (mean $\pm$ std over 3 seeds). Reference numbers are from Cohen et al. [2018]. σ_{rot}: standard deviation of accuracy under 10 random $\mathrm{SO}(3)$ rotations (NR-trained models).

Model	Params	NR/NR ↑	R/R ↑	NR/R ↑	σ_{rot} ↓
Cohen planar CNN	—	0.98	0.23	0.11	—
Cohen spherical CNN	—	0.96	0.95	0.94	—
Standard CNN	185K	0.991±0.000	0.781±0.003	0.138±0.003	0.044
Power Spec. + MLP	11K	0.780±0.002	0.780±0.002	0.777±0.001	3×10^{-4}
Power Spec. + MLP (matched)	166K	0.794±0.011	0.794±0.011	0.792±0.010	3×10^{-4}
`SO3onS2` **Bispectrum + MLP (ours)**	232K	**0.950±0.000**	**0.949±0.002**	**0.951±0.001**	$\mathbf{1\times10^{-4}}$

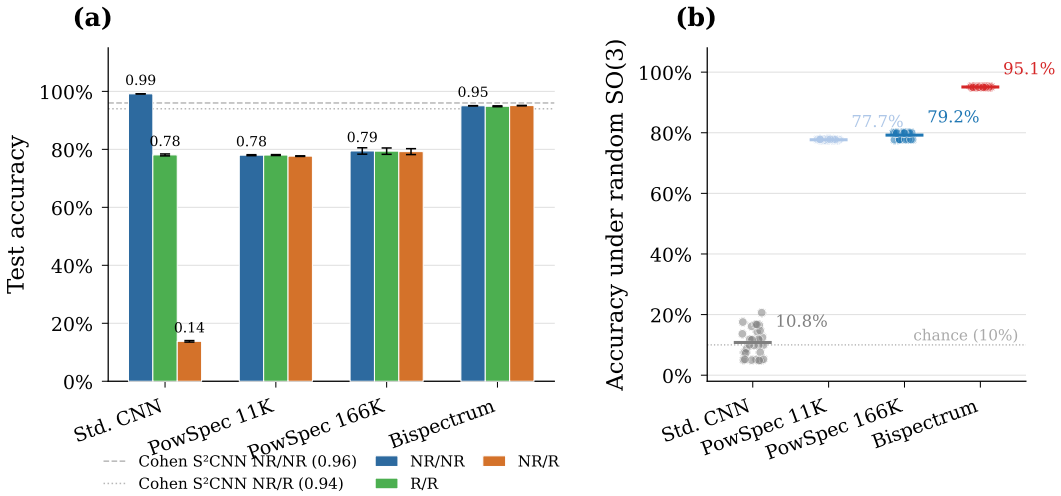

Figure 4: Spherical MNIST analysis (`SO3onS2` selective bispectrum). **(a)** Cohen et al. protocol: NR/NR, R/R, and NR/R accuracy. Dashed lines show Cohen's spherical CNN reference. **(b)** Accuracy under 10 random $\mathrm{SO}(3)$ rotations (NR-trained models, 30 points per model = 10 rotations $\times$ 3 seeds). Scaling the power spectrum MLP from 11K to 166K parameters yields only +1.4pp; the `SO3onS2` bispectrum at 232K reaches 95.0%.

Completeness gap. The power spectrum is equally invariant but reaches only 78%—a 17pp deficit that persists even after scaling the MLP to 166K parameters (+1.4pp), confirming incompleteness rather than capacity as the bottleneck.

Comparison with Cohen et al. Cohen's spherical CNN [Cohen et al., 2018] reaches 96%/95%/94% with learned equivariant convolutions. Our single-pass bispectrum (95.0%/94.9%/95.1%) is 1–2pp lower but has a tighter invariance gap (0.1pp vs. 2.0pp).

6 Selective $\mathrm{SO}(3)$ bispectrum on S^2

6.1 Setup

A real signal $f : S^2 \to \mathbb{R}$ band-limited at degree L expands in spherical harmonics as $f(\theta,\varphi) = \sum_{\ell=0}^{L}\sum_{m=-\ell}^{\ell} a_\ell^m Y_\ell^m(\theta,\varphi)$ with the reality constraint $a_\ell^{-m} = (-1)^m \overline{a_\ell^m}$, so it is parameterised by $(L{+}1)^2$ real scalars (Definition 5, eq. (5)). On $S^2 = \mathrm{SO}(3)/\mathrm{SO}(2)_z$, the relevant irreducible representations of $\mathrm{SO}(3)$ are $\{\rho_\ell\}_{\ell\geq 0}$ with V_{ρ_ℓ} the $(2\ell{+}1)$-dimensional space of degree-ℓ spherical harmonics and Wigner-D action $\rho_\ell(g) = D^\ell(g)$. In the irrep notation of Mataigne et al. [2024], the Fourier coefficient at ρ_ℓ is $\mathcal{F}(f)_{\rho_\ell} = \mathbf{F}_\ell \in \mathbb{C}^{2\ell+1}$; we write $\mathbf{F}_\ell$ throughout. The scalar ρ_ℓ-isotypic

component of the matrix-valued bispectrum $\beta_{\rho_{\ell_1},\rho_{\ell_2}}(f)$ is the familiar

$$\beta_{\ell_1,\ell_2,\ell} = \sum_{\substack{m_1,m_2,m \\ m_1+m_2=m}} C^{\ell,\,m}_{\ell_1,\,m_1;\,\ell_2,\,m_2}\, a^{m_1}_{\ell_1}\, a^{m_2}_{\ell_2}\, \overline{a^{m}_{\ell}}, \tag{3}$$

the specialisation of eq. (2) to SO(3)-on-S^2. The full bispectrum has $\Theta(L^3)$ such scalars; the orbit space $\mathbb{R}^{(L+1)^2}/\mathrm{SO}(3)$ has dimension $(L+1)^2 - 3$, so any complete invariant requires $\Omega(L^2)$ independent components (Proposition 9). The Mataigne chain $\beta_{1,\ell,\ell+1}$ for $\ell = 0, \ldots, L-1$ provides only $\sim 3L$ scalars and cannot achieve this bound.

6.2 Selective construction and algorithm

Let $L_{\mathrm{seed}} := \min(L, 7)$. The *seed index set* collects all admissible bispectral triples up to L_{seed},

$$\mathcal{S}_{\mathrm{seed}} := \{(\ell_1, \ell_2, \ell) : 0 \le \ell_1 \le \ell_2 \le L_{\mathrm{seed}},\ |\ell_1 - \ell_2| \le \ell \le \min(\ell_1 + \ell_2, L_{\mathrm{seed}})\},$$

of constant size ≤ 512 once $L \ge 7$. For each high-degree target $\ell \ge 8$, we use the size-$(2\ell+1)$ *bootstrap block* $\mathcal{T}_\ell$ of eq. (8) (Definition 11), comprising the entries linear in $\mathbf{F}_\ell$ with coefficients depending only on already-recovered lower degrees. The *selective* SO(3) *bispectrum on* S^2 is then the concatenation

$$\beta_{\mathrm{sel}}(f) := \left\{\beta_{\ell_1,\ell_2,\ell}\right\}_{(\ell_1,\ell_2,\ell)\in\mathcal{S}_\beta}, \qquad \mathcal{S}_\beta := \mathcal{S}_{\mathrm{seed}} \cup \bigcup_{\ell=8}^{L} \mathcal{T}_\ell, \tag{4}$$

storing each entry's full complex value (real and imaginary parts). Algorithm 1 writes this as a Mataigne-style BFS over selected triples; the implementation-grade index set (Algorithm 2 in Appendix A) augments $\mathcal{S}_\beta$ with mandatory even self-couplings and optional CG-power scalars for numerical robustness, and is a strict superset of β_{sel}.

Algorithm 1 Selective SO(3) bispectrum on S^2 (BFS form, mirroring Mataigne et al. [2024], Alg. 1).

```
Input: band-limit L, seed cutoff L_seed = min(L, 7)
L_β ← all admissible triples (ℓ1, ℓ2, ℓ) with ℓ ≤ L_seed {seed block S_seed}
for ℓ = L_seed+1, ..., L do
    append the bootstrap block T_ℓ to L_β {2ℓ+1 entries linear in F_ℓ}
end for
return β_sel(f) = {β_{ℓ1,ℓ2,ℓ}(f) : (ℓ1, ℓ2, ℓ) ∈ L_β}
```

The output size is $|\beta_{\mathrm{sel}}| = |\mathcal{S}_{\mathrm{seed}}| + \sum_{\ell=8}^{L}(2\ell+1) = \Theta(L^2)$, matching the orbit-space lower bound up to a constant (Proposition 13).

Remark 1 (CG power: optional implementation diagnostic). The library appends the optional CG power spectrum $P_{\ell_1,\ell_2,\ell} := \|(\mathbf{F}_{\ell_1} \otimes \mathbf{F}_{\ell_2})|_\ell\|^2$ (Definition 6) at $O(L^2)$ extra cost; under the hypotheses below the CG-power augmentation is unnecessary for completeness, and we set $\mathcal{S}_P = \varnothing$ in the formal theorem. The implementation includes it because (i) it provides a per-degree numerical robustness diagnostic, and (ii) it costs $O(L^2)$ extra scalars without changing the asymptotic budget.

6.3 Conditional completeness theorem

Theorem 2 reduces completeness of β_{sel} to two named hypotheses; both are stated formally in Appendix A.

H1 (finite seed injectivity, Hypothesis 24). The full scalar bispectrum at $L_{\mathrm{seed}} \le 7$ is injective on a Zariski-open gauge-fixed stratum of $V_{L_{\mathrm{seed}}}$, with T_R resolved by the odd-parity entry $\mathrm{Im}(\beta_{2,3,4})$. Structural support: Kakarala Theorem 7 [Kakarala, 1992, 2012] for $L^1(\mathrm{SO}(3))$, plus the exact-arithmetic Gröbner-basis certificates of Bandeira et al. [2017] for $L_0 \le 15$; neither alone is a band-limited proof.

H2 (uniform bootstrap-determinant nonvanishing, Hypothesis 32). For every $\ell \ge 8$, $\Delta_\ell := \det A_\ell$ is not identically zero as a polynomial on the lower-degree real coordinates, where $A_\ell \in \mathbb{R}^{(2\ell+1)\times(2\ell+1)}$ is the bootstrap coefficient matrix assembled from $\mathcal{T}_\ell$ on the real-signal slice.

Theorem 2 (Conditional completeness of β_{sel})**.** *Let $L \geq 4$ and assume H1 and, if $L \geq 8$, H2. For $f, f' : S^2 \to \mathbb{R}$ band-limited at degree L and lying outside a proper algebraic subvariety of V_L (Assumption 15), if $\beta_{\mathrm{sel}}(f) = \beta_{\mathrm{sel}}(f')$, then $f' = g \cdot f$ for some $g \in \mathrm{SO}(3)$. The output β_{sel} has $\Theta(L^2)$ real components, matching the orbit-space dimension $(L+1)^2 - 3$ up to a constant factor.*

Inversion sketch (four steps). (i) *Gauge-fix* $\mathbf{F}_1$ along $\hat{z}$ and the residual $\mathrm{SO}(2)_z$ phase via $\arg a_2^1$, consuming all three rotational DOFs (Lemma 18). (ii) *Closed form for* $\mathbf{F}_0, \mathbf{F}_1$: $\beta_{0,0,0} = (a_0^0)^3$ pins A_0, then $\beta_{0,1,1} = A_0 \|\mathbf{F}_1\|^2$ pins $\mathbf{F}_1 = (0, c, 0)^\top$ (Lemma 20). (iii) *Seed solve at* $\ell \leq L_{\mathrm{seed}}$: under H1, the seed block $\mathcal{S}_{\mathrm{seed}}$ uniquely determines $(\mathbf{F}_2, \ldots, \mathbf{F}_{L_{\mathrm{seed}}})$ on the gauge slice; the parity entry $\mathrm{Im}(\beta_{2,3,4})$ separates f from $T_R \cdot f$ (Corollary 30; on the "up-to-reflection" caveat of Edidin and Satriano [2024] see Remark 25). (iv) *Linear bootstrap for* $\ell \geq 8$: under H2, the system $A_\ell \, \mathbf{F}_\ell = b_\ell$ is invertible on a Zariski-open set of lower-degree coordinates, and induction recovers $\mathbf{F}_8, \ldots, \mathbf{F}_L$.

6.4 Computational evidence and scope

Hypothesis H2 is verified per degree at the deterministic seed-42 witness in IEEE 754 double precision: for every $\ell \in \{8, \ldots, 100\}$, the singular-value ratio $\sigma_{\min}(A_\ell)/\sigma_{\max}(A_\ell)$ exceeds 10^{-4} a margin of more than 10^{11} above double-precision rounding error. Scripts: `paper/exact_certificates.py` and `benchmarks/verify_linear_bootstrap.py`. The structural form of $\mathcal{T}_\ell$ is uniform in $\ell \geq 8$, so extending the verified range beyond $L = 100$ requires only re-running the script with a larger upper bound. Hypothesis H1 is supported by Kakarala Theorem 7 [Kakarala, 2012, 1992] as structural motivation and by the exact-arithmetic certificates of Bandeira et al. [2017] for $L_0 \leq 15$, but neither alone constitutes a band-limited proof of H1. Three routes would upgrade H2 to an unconditional theorem: interval arithmetic at the witness, exact rational determinants per ℓ, or a closed-form Wigner-3j non-vanishing identity. Full proofs and the complete set of hypotheses are in Appendix A.

7 Discussion and Conclusion

We have presented `bispectrum`, a PyTorch library that makes selective G-bispectral invariants practical for machine learning. The library covers seven group/domain pairs with a uniform `nn.Module` interface and required one technical contribution during implementation: a selective $\mathrm{SO}(3)$ bispectrum on S^2 achieving the optimal $\Theta(L^2)$ coefficient count (Theorem 2). Table 2 confirms that the selective forward pass adds negligible overhead to training pipelines (<1 ms for all finite-group modules on a single GPU). Across three experimental settings, bispectral pooling provides numerically exact rotation invariance and a data-efficiency advantage over incomplete alternatives in the low-data / moderate-capacity regime, where the complete invariant map provides a stronger inductive bias that partially compensates for limited training data.

When not to use the bispectrum. The data-efficiency advantage of complete invariants is most pronounced in the low-data, moderate-capacity regime. At high model capacity (Section 5), incomplete alternatives such as max pooling can match or surpass bispectral pooling, likely because the network learns to compensate for missing invariant information. Each bispectrum module is also invariant to a *specific* group (e.g. rotations only, or translations only); for tasks that need invariance to multiple or unknown symmetries, learned invariance via data augmentation or group-equivariant networks with flexible symmetry may be a better fit. Completeness is a property of the *map*, not of downstream task performance: a complete invariant provides a stronger inductive bias, but does not guarantee superior accuracy on every task. Invertibility of the bispectrum for generic signals was first proved by Giannakis [1989]; Kakarala [2012] extended this to compact groups. For $\mathrm{SO}(3)$ on S^2, computational cost grows as $O(L^2)$ per forward pass; at very large band-limits this may become a bottleneck relative to simpler pooling strategies.

Limitations. Inversion for $\mathrm{SO}(3)$ on S^2 is not yet implemented; only the selective forward map is available. The selective $\mathrm{SO}(3)$ completeness statement is *conditional*: it depends on Hypothesis H1 (finite-band seed injectivity, motivated by Kakarala Theorem 7 [Kakarala, 1992, 2012] and the Gröbner-basis evidence of Bandeira et al. [2017] for $L_0 \leq 15$) and Hypothesis H2 (uniform bootstrap-determinant nonvanishing). H2 is supported per ℓ at the deterministic seed-42 witness with $\sigma_{\min}/\sigma_{\max} \geq 10^{-4}$ in double precision (margin $>10^{11}$ above rounding error) for every $\ell \leq 100$;

the constant-size seed Jacobian rank (Lemma 39) is verified in exact rational arithmetic via `sympy`. Upgrading H2 to an unconditional theorem requires interval arithmetic, exact rational determinants per ℓ, or a closed-form Wigner-3j non-vanishing identity; a full characterisation of the degenerate set is open. Our experimental comparisons do not include tensor-product architectures (e3nn, MACE) or wavelet scattering, which target different input assumptions and would require non-trivial adaptation to the benchmarks used here.

Future work. Key directions include: implementing $\mathrm{SO}(3)$ inversion from selective bispectral coefficients; a structural proof of bootstrap rank for all L; numerical conditioning analysis at large bandlimits; broader framework integration (the library is intentionally framework-agnostic and already drops in as an invariant module for architectures built on `escnn` [Cesa et al., 2022] or `e3nn` [Geiger and Smidt, 2022]); and applications to cryo-EM reconstruction and molecular property prediction.

References

Afonso S. Bandeira, Ben Blum-Smith, Joe Kileel, Amelia Perry, Jonathan Weed, and Alexander S. Wein. Estimation under group actions: recovering orbits from invariants. *arXiv preprint arXiv:1712.10163*, 2017. URL https://arxiv.org/abs/1712.10163.

Ilyes Batatia, Gábor Csányi, Dávid P. Kovács, Christoph Ortner, and Gregor Simm. MACE: Higher order equivariant message passing neural networks for fast and accurate force fields. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 11423–11436, 2022. doi: 10.52202/068431-0830. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/4a36c3c51af11ed9f34615b81edb5bbc-Abstract-Conference.html.

Ilyes Batatia, Mario Geiger, Jose Munoz, Tess Smidt, Lior Silberman, and Christoph Ortner. A general framework for equivariant neural networks on reductive Lie groups, 2023. URL https://arxiv.org/abs/2306.00091.

Tamir Bendory, Dan Edidin, Josh Katz, and Shay Kreymer. Orbit recovery for spherical functions. *arXiv preprint arXiv:2508.02674*, 2025. URL https://arxiv.org/abs/2508.02674.

Thomas S. Binns, Franziska Pellegrini, Tin Jurhar, Tien D. Nguyen, Richard M. Köhler, and Stefan Haufe. PyBispectra: A toolbox for advanced electrophysiological signal processing using the bispectrum. *Journal of Open Source Software*, 10(112):8504, 2025. doi: 10.21105/joss.08504. URL https://doi.org/10.21105/joss.08504.

Boris Bonev, Thorsten Kurth, Christian Hundt, Jaideep Pathak, Maximilian Balu, Karthik Kashinath, and Animashree Anandkumar. Spherical Fourier neural operators: Learning stable dynamics on the sphere. In *International Conference on Machine Learning (ICML)*, 2023.

Gabriele Cesa, Leon Lang, and Maurice Weiler. A program to build E(N)-equivariant steerable CNNs. In *International Conference on Learning Representations (ICLR)*, 2022. URL https://openreview.net/forum?id=WE4qe9xlnQw.

Arthur Chevalley, Valentin Oreiller, Julien Fageot, John O. Prior, Vincent Andrearczyk, and Adrien Depeursinge. A bispectral 3D U-Net for rotation robustness in medical segmentation. In *Topology- and Graph-Informed Imaging Informatics*, pages 43–54. Springer Nature Switzerland, 2024. doi: 10.1007/978-3-031-73967-5_5. URL https://doi.org/10.1007/978-3-031-73967-5_5.

Taco S. Cohen and Max Welling. Group equivariant convolutional networks. In *International Conference on Machine Learning (ICML)*, pages 2990–2999, 2016.

Taco S. Cohen, Mario Geiger, Jonas Köhler, and Max Welling. Spherical CNNs. In *International Conference on Learning Representations (ICLR)*, 2018.

Persi Diaconis and Daniel Rockmore. Efficient computation of the Fourier transform on finite groups. *Journal of the American Mathematical Society*, 3(2):297–332, April 1990. doi: 10.1090/S0894-0347-1990-1030655-4. URL https://www.ams.org/journals/jams/1990-03-02/S0894-0347-1990-1030655-4/.

Dan Edidin and Matthew Satriano. Orbit recovery for band-limited functions. *SIAM Journal on Applied Algebra and Geometry*, 8(3):733–755, 2024. doi: 10.1137/23M1577808. URL https://arxiv.org/abs/2306.00155.

Babak Ehteshami Bejnordi, Mitko Veta, Paul Johannes van Diest, Bram van Ginneken, Nico Karssemeijer, Geert Litjens, Jeroen A. W. M. van der Laak, and the CAMELYON16 Consortium. Diagnostic assessment of deep learning algorithms for detection of lymph node metastases in women with breast cancer. *JAMA*, 318(22):2199–2210, 2017. doi: 10.1001/jama.2017.14585.

Carlos Esteves, Christine Allen-Blanchette, Ameesh Makadia, and Kostas Daniilidis. Learning SO(3) equivariant representations with spherical CNNs. In *European Conference on Computer Vision (ECCV)*, pages 52–68. Springer, 2018.

Mario Geiger and Tess Smidt. e3nn: Euclidean neural networks. *arXiv preprint arXiv:2207.09453*, 2022. URL https://arxiv.org/abs/2207.09453.

Georgios B. Giannakis. Signal reconstruction from multiple correlations: frequency- and time-domain approaches. *Journal of the Optical Society of America A*, 6(5):682–697, May 1989. doi: 10.1364/JOSAA.6.000682. URL https://opg.optica.org/josaa/abstract.cfm?URI=josaa-6-5-682.

Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, et al. Array programming with NumPy. *Nature*, 585(7825):357–362, 2020.

Gao Huang, Zhuang Liu, Laurens van der Maaten, and Kilian Q. Weinberger. Densely connected convolutional networks. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.

Ramakrishna Kakarala. *Triple Correlation on Groups*. PhD thesis, University of California, Irvine, 1992.

Ramakrishna Kakarala. The bispectrum as a source of phase-sensitive invariants for Fourier descriptors: A group-theoretic approach. *Journal of Mathematical Imaging and Vision*, 44(3):341–353, 2012. doi: 10.1007/s10851-012-0330-6. URL https://arxiv.org/abs/0902.0196.

Thijs P. Kuipers and Erik J. Bekkers. Regular SE(3) group convolutions for volumetric medical image analysis. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2023*, pages 252–261. Springer Nature Switzerland, 2023. doi: 10.1007/978-3-031-43898-1_25. URL https://doi.org/10.1007/978-3-031-43898-1_25.

Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019. URL https://openreview.net/forum?id=Bkg6RiCqY7.

Stéphane Mallat. Group invariant scattering. *Communications on Pure and Applied Mathematics*, 65(10):1331–1398, 2012. doi: 10.1002/cpa.21413.

Simon Mataigne, Johan Mathe, Sophia Sanborn, Christopher Hillar, and Nina Miolane. The selective G-bispectrum and its inversion: Applications to G-invariant networks. In *Advances in Neural Information Processing Systems*, volume 37, pages 115682–115711. Neural Information Processing Systems Foundation, Inc., 2024. doi: 10.52202/079017-3674. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/d1a1e8713fcd5626656553c82f7c3b26-Abstract-Conference.html.

Sophia Myers and Nina Miolane. The selective disk bispectrum and its inversion, with application to multi-reference alignment. *arXiv preprint arXiv:2511.19706*, 2025. URL https://arxiv.org/abs/2511.19706.

Valentin Oreiller, Julien Fageot, Vincent Andrearczyk, John O. Prior, and Adrien Depeursinge. Robust multi-organ nucleus segmentation using a locally rotation invariant bispectral U-Net. In *Proceedings of The 5th International Conference on Medical Imaging with Deep Learning*, volume 172 of *Proceedings of Machine Learning Research*, pages 929–943. PMLR, 2022. URL https://proceedings.mlr.press/v172/oreiller22a.html.

Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zach DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

Sophia Sanborn, Christian Shewmake, Bruno Olshausen, and Christopher Hillar. Bispectral neural networks. In *International Conference on Learning Representations (ICLR)*, 2023.

Mateus Sangalli, Samy Blusseau, Santiago Velasco-Forero, and Jesús Angulo. Moving frame net: SE(3)-equivariant network for volumes. In *Proceedings of the 1st NeurIPS Workshop on Symmetry and Geometry in Neural Representations*, volume 197 of *Proceedings of Machine Learning Research*, pages 81–97. PMLR, 2023. URL `https://proceedings.mlr.press/v197/sangalli23a.html`.

Nathaniel Thomas, Tess Smidt, Steven Kearnes, Lusann Yang, Li Li, Kai Kohlhoff, and Patrick Riley. Tensor field networks: Rotation- and translation-equivariant neural networks for 3D point clouds. *arXiv preprint arXiv:1802.08219*, 2018. URL `https://arxiv.org/abs/1802.08219`.

Bastiaan S. Veeling, Jasper Linmans, Jim Winkens, Taco Cohen, and Max Welling. Rotation equivariant CNNs for digital pathology. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, pages 210–218. Springer, 2018.

Maurice Weiler and Gabriele Cesa. General E(2)-equivariant steerable CNNs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

YuQing Xie, Ameya Daigavane, Mit Kotak, and Tess Smidt. The price of freedom: Exploring expressivity and runtime tradeoffs in equivariant tensor products. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, volume 267 of *Proceedings of Machine Learning Research*, pages 68599–68625. PMLR, 2025.

Jiancheng Yang, Rui Shi, Donglai Wei, Zequan Liu, Lin Zhao, Bilian Ke, Hanspeter Pfister, and Bingbing Ni. MedMNIST v2 – a large-scale lightweight benchmark for 2D and 3D biomedical image classification. *Scientific Data*, 10(1):41, 2023.

Dmitrii Zhemchuzhnikov and Sergei Grudinin. ILPO-NET: Network for the invariant recognition of arbitrary volumetric patterns in 3D. In *Machine Learning and Knowledge Discovery in Databases. Research Track*, pages 352–368. Springer Nature Switzerland, 2024. doi: 10.1007/978-3-031-70359-1_21. URL `https://doi.org/10.1007/978-3-031-70359-1_21`.

Dmitrii Zhemchuzhnikov and Sergei Grudinin. On the Fourier analysis in the SO(3) space: the EquiLoPO network. In *International Conference on Learning Representations (ICLR)*, 2025. URL `https://openreview.net/forum?id=LvTSvdiSwG`.

A Selective $\mathrm{SO}(3)$ Bispectrum – Full Proof

This appendix proves the conditional completeness theorem of Section 6. The two load-bearing assumptions H1 (finite seed injectivity) and H2 (uniform bootstrap-determinant nonvanishing) named in the main text are stated formally below as Hypotheses 24 and 32. The proof follows the Kakarala pattern [Kakarala, 1992, 2012]: a constant-size full-bispectrum seed at degree $L_{\mathrm{seed}} \leq 7$ followed by a triangular per-degree linear recurrence in which each $\mathbf{F}_\ell$ ($\ell \geq 8$) is recovered by solving an explicit $(2\ell{+}1)$-dimensional system. Clebsch–Gordan power entries are optional implementation diagnostics and are not used in the proof ($\mathcal{S}_P = \varnothing$ throughout). The implementation-grade selective-index-set construction (Algorithm 2) is preserved at the end of Section A.3.

The total invariant size is therefore $|\mathcal{S}_{\mathrm{seed}}| + \sum_{\ell=8}^{L}(2\ell{+}1) = \Theta(L^2)$, matching the orbit-space lower bound $(L{+}1)^2 - 3$ of Proposition 9 up to a constant factor (Proposition 13).

Table 6: Asymptotic entry counts in β_{sel}. The constant-size seed block $\mathcal{S}_{\mathrm{seed}}$ collects every triple at degrees $\leq L_{\mathrm{seed}} = \min(L, 7)$; the high-degree selective part allocates exactly $2\ell{+}1$ rows per target degree (the bootstrap block $\mathcal{T}_\ell$ of eq. (8)).

degree range	selected β	total
$\ell \leq 7$	full seed, ≤ 512	$O(1)$
$\ell \geq 8$	$2\ell{+}1$	$O(\ell)$

A.1 Setup and notation

Definition 3 (Band-limited signal on S^2)**.** A signal $f : S^2 \to \mathbb{R}$ is *band-limited* at degree L if its spherical harmonic expansion

$$f(\theta, \varphi) = \sum_{\ell=0}^{L} \sum_{m=-\ell}^{\ell} a_\ell^m \, Y_\ell^m(\theta, \varphi) \tag{5}$$

has $a_\ell^m = 0$ for all $\ell > L$. We write $\mathbf{F}_\ell = (a_\ell^{-\ell}, \ldots, a_\ell^\ell) \in \mathbb{C}^{2\ell+1}$ for the coefficient vector at degree ℓ.

For real f, the conjugation symmetry $a_\ell^{-m} = (-1)^m \overline{a_\ell^m}$ holds, so $\mathbf{F}_\ell$ has $2\ell + 1$ real degrees of freedom. The total signal dimension is $\dim = \sum_{\ell=0}^{L}(2\ell + 1) = (L + 1)^2$.

Definition 4 (Rotation action)**.** Under $g \in \mathrm{SO}(3)$, the spherical harmonic coefficients transform as $\mathbf{F}_\ell \mapsto D^\ell(g)\, \mathbf{F}_\ell$, where D^ℓ is the $(2\ell + 1)$-dimensional Wigner D-matrix.

In the irrep notation of Mataigne et al. [2024], $\{\rho_\ell\}_{\ell \geq 0}$ are the (unitary) irreducible representations of $\mathrm{SO}(3)$ with V_{ρ_ℓ} the $(2\ell{+}1)$-dimensional space of degree-ℓ spherical harmonics and $\rho_\ell(g) = D^\ell(g)$; the Fourier coefficient of f at ρ_ℓ is $\mathcal{F}(f)_{\rho_\ell} = \mathbf{F}_\ell$. We use $\mathbf{F}_\ell$ as the working symbol throughout.

Definition 5 (Bispectrum on S^2)**.** For each triple (ℓ_1, ℓ_2, ℓ) satisfying the triangle inequality $|\ell_1 - \ell_2| \leq \ell \leq \ell_1 + \ell_2$, the *bispectral coefficient* is

$$\beta_{\ell_1,\ell_2,\ell} = \sum_{\substack{m_1,\, m_2,\, m \\ m_1+m_2=m}} C^{\ell,\, m}_{\ell_1,\, m_1;\, \ell_2,\, m_2} \; a_{\ell_1}^{m_1} \, a_{\ell_2}^{m_2} \, \overline{a_\ell^m}, \tag{6}$$

where $C^{\ell,\, m}_{\ell_1,\, m_1;\, \ell_2,\, m_2}$ denotes the Clebsch–Gordan coefficient $\langle \ell_1, m_1;\, \ell_2, m_2 \mid \ell, m \rangle$. Each $\beta_{\ell_1,\ell_2,\ell}$ is a scalar; for real f it is real when $\ell_1{+}\ell_2{+}\ell$ is even and purely imaginary when $\ell_1{+}\ell_2{+}\ell$ is odd (Lemma 27).

Definition 5 is the specialization of the general G-bispectrum (eq. (2)) to the case $G = \mathrm{SO}(3)$ acting on S^2, where the irreps are Wigner-D matrices and the Fourier coefficients are spherical harmonic coefficients a_ℓ^m.

Definition 6 (CG power spectrum, optional diagnostic)**.** For a triple (ℓ_1, ℓ_2, ℓ) satisfying the triangle inequality, write $(\mathbf{F}_{\ell_1} \otimes \mathbf{F}_{\ell_2})\big|_\ell$ for the projection of the Kronecker product $\mathbf{F}_{\ell_1} \otimes \mathbf{F}_{\ell_2}$ onto the D^ℓ-isotypic component, realised by contracting with the CG coefficients $C^{\ell,\, m}_{\ell_1,\, m_1;\, \ell_2,\, m_2}$. The *CG power spectrum entry* is the squared norm of this projection,

$$P_{\ell_1,\ell_2,\ell} := \|(\mathbf{F}_{\ell_1} \otimes \mathbf{F}_{\ell_2})|_\ell\|^2 = \sum_{m=-\ell}^{\ell} \Big| \sum_{\substack{m_1, m_2 \\ m_1+m_2=m}} C^{\ell,\, m}_{\ell_1,\, m_1;\, \ell_2,\, m_2} \; a_{\ell_1}^{m_1} \, a_{\ell_2}^{m_2} \Big|^2. \tag{7}$$

This is a degree-4, real-valued, $\mathrm{SO}(3)$-invariant polynomial in the signal coefficients. CG power entries are an *optional* implementation diagnostic and are not used in the conditional completeness theorem; we set $\mathcal{S}_P = \varnothing$ throughout the formal proof. The library may append a $O(L^2)$ family of CG power scalars without affecting completeness or the asymptotic output size.

Proposition 7 (Size of the full bispectrum)**.** *The number of admissible triples* (ℓ_1, ℓ_2, ℓ) *with* $0 \leq \ell_1, \ell_2, \ell \leq L$ *satisfying the triangle inequality* $|\ell_1 - \ell_2| \leq \ell \leq \ell_1 + \ell_2$ *is* $\Theta(L^3)$.

Proof. For fixed ℓ_1 and ℓ_2, the triangle inequality admits $\ell \in \{|\ell_1 - \ell_2|, \ldots, \min(\ell_1 + \ell_2, L)\}$, giving $\min(\ell_1 + \ell_2, L) - |\ell_1 - \ell_2| + 1$ valid values of ℓ. Summing:

$$N_{\text{full}} = \sum_{\ell_1=0}^{L} \sum_{\ell_2=0}^{L} \big(\min(\ell_1 + \ell_2, L) - |\ell_1 - \ell_2| + 1\big) = \tfrac{1}{3}L^3 + \mathcal{O}(L^2).$$

□

Remark 8 (Invariance). Each $\beta_{\ell_1,\ell_2,\ell}$ is $\mathrm{SO}(3)$-invariant: the Wigner matrices satisfy $D^{\ell_1} \otimes D^{\ell_2} = \bigoplus_\ell D^\ell$ under the CG decomposition, so the triple product contracted with CG coefficients is unchanged by $\mathbf{F}_\ell \mapsto D^\ell(g)\mathbf{F}_\ell$.

A.2 Why $\mathcal{O}(L)$ is impossible

Proposition 9 (Dimensional lower bound)**.** *Any complete* $\mathrm{SO}(3)$*-invariant for band-limited signals on* S^2 *at band-limit* L *must have at least* $(L+1)^2 - 3$ *functionally independent components.*

Proof. The orbit space $\mathbb{R}^{(L+1)^2}/\mathrm{SO}(3)$ has dimension $(L+1)^2 - \dim \mathrm{SO}(3) = (L+1)^2 - 3$. Any injective map from the orbit space to $\mathbb{R}^k$ requires $k \geq (L+1)^2 - 3$, hence $\Omega(L^2)$. □

Remark 10. The chain $\beta_{1,\ell,\ell+1}$ for $\ell = 0, \ldots, L-1$ provides only $\sim 3L$ scalar constraints (since each chain entry couples $\mathbf{F}_1 \in \mathbb{C}^3$ with one known $\mathbf{F}_\ell$). This cannot recover $(L+1)^2 \approx L^2$ unknowns.

A.3 The selective index set

The goal is to select, for each target degree $\ell \geq 8$, exactly $2\ell{+}1$ bispectral triples (ℓ_1, ℓ_2, ℓ) such that the selected entries are simultaneously linear in $\mathbf{F}_\ell$ (or in $\overline{\mathbf{F}_\ell}$, which we conjugate) and have a coefficient matrix depending only on already-recovered lower degrees. Combined with a constant-size full-bispectrum seed at degrees $\leq L_{\text{seed}} = \min(L, 7)$, this gives an $O(L^2)$-sized invariant matching the orbit-space lower bound of Proposition 9.

Definition 11 (Selective $\mathrm{SO}(3)$ bispectrum on S^2)**.** Set $L_{\text{seed}} := \min(L, 7)$. The *seed index set* is the full bispectrum truncated at L_{seed},

$$\mathcal{S}_{\text{seed}} := \{(\ell_1, \ell_2, \ell) : 0 \leq \ell_1 \leq \ell_2 \leq L_{\text{seed}},\ |\ell_1{-}\ell_2| \leq \ell \leq \min(\ell_1{+}\ell_2, L_{\text{seed}})\},$$

of size at most $(L_{\text{seed}}{+}1)^3 \leq 512$, independent of L once $L \geq 7$. For each high-degree target $\ell \geq 8$, the *bootstrap block* $\mathcal{T}_\ell$ (also used in Proposition 12) is

$$\begin{aligned} \mathcal{T}_\ell := &\{(a, \ell, \ell - a) : 1 \leq a \leq \ell{-}1\} \\ &\cup \{(a, \ell, \ell - a + 1) : 2 \leq a \leq \ell{-}1\} \\ &\cup \{(a, \ell - a, \ell) : 1 \leq a \leq 4\} \end{aligned} \tag{8}$$

with $(\ell{-}1) + (\ell{-}2) + 4 = 2\ell{+}1$ entries. The bispectral index set is

$$\mathcal{S}_\beta := \mathcal{S}_{\text{seed}} \cup \bigcup_{\ell=8}^{L} \mathcal{T}_\ell.$$

The *selective bispectrum* $\beta_{\text{sel}} : V_L \to \mathbb{R}^N$ is the concatenation of all selected scalar entries $\{\beta_{\ell_1,\ell_2,\ell}\}_{(\ell_1,\ell_2,\ell)\in\mathcal{S}_\beta}$, storing the full complex value of each entry (equivalently, real and imaginary parts). We set $\mathcal{S}_P := \varnothing$ in the formal theorem; the optional CG power augmentation (Definition 6) may be appended at $O(L^2)$ cost without changing the asymptotic budget or the conclusion of Theorem 17.

Library implementation note. Algorithm 2 below specifies the implementation-grade index construction used by the `SO3onS2` module. It augments Definition 11 in two ways: (i) it includes *mandatory even self-couplings* $\beta_{\ell,\ell,\ell'}$ (ℓ' even, $2 \leq \ell' \leq \ell$) and the power entries $\beta_{0,\ell,\ell}$ for numerical robustness on the real-signal slice, and (ii) it greedily appends CG-power entries to $\mathcal{S}_P$ to maximise the per-degree real Jacobian rank. The augmented invariant is a strict superset of β_{sel}, so completeness transfers automatically; the formal theorem is stated for the smaller, purely-bispectral β_{sel}.

Algorithm 2 Augmented selective index set for the $\mathrm{SO}(3)$ bispectrum on S^2 at band-limit L. Returns bispectral triples $\mathcal{S}_\beta$ and CG power triples $\mathcal{S}_P$. This extends the BFS-based selective index construction of Mataigne et al. [2024] from finite groups to $\mathrm{SO}(3)$ on S^2, adding the CG power augmentation step that is specific to real signals.

Input: Band-limit L
Output: Bispectral set $\mathcal{S}_\beta$, CG power set $\mathcal{S}_P$
$\mathcal{S}_\beta \leftarrow \{(0,0,0)\}$; $\mathcal{S}_P \leftarrow \emptyset$
Seed block: add to $\mathcal{S}_\beta$ all triples (ℓ_1, ℓ_2, ℓ) with $0 \le \ell_1 \le \ell_2 \le L_0 = 4$ and $|\ell_1 - \ell_2| \le \ell \le \min(\ell_1 + \ell_2, L_0)$ {constant-size full-bispectrum seed feeding Hypothesis 24}
for $\ell = 1, \ldots, L$ **do**
  budget $\leftarrow 2\ell + 1$
  if $\ell = 4$ **then**
    budget $\leftarrow 10$ {keep all chain+cross at $\ell = 4$}
  end if
  candidates $\leftarrow [\,]$
  if $4 \le \ell \le 7$ **then**
    append the explicit low-degree linear block $\mathcal{T}_\ell^{\mathrm{small}}$ to candidates
  else if $\ell \ge 8$ **then**
    for $a = 1, \ldots, \ell{-}1$ **do**
      append $(a,\ \ell,\ \ell{-}a)$ to candidates
    end for
    for $a = 2, \ldots, \ell{-}1$ **do**
      append $(a,\ \ell,\ \ell{-}a{+}1)$ to candidates
    end for
    for $a = 1, 2, 3, 4$ **do**
      append $(a,\ \ell{-}a,\ \ell)$ to candidates
    end for
  else
    append low-degree chain candidates
    if $\ell \ne 2$ **then**
      append cross candidates {$\beta_{1,2,1}$ excluded}
    end if
  end if
  append $(0,\ \ell,\ \ell)$ to candidates {power entry}
  if $\ell \notin \{1,\ 2\}$ **then**
    for $\ell' = \ell{-}1, \ldots, 0$ **do**
      append $(\ell,\ \ell,\ \ell')$ to candidates
    end for
  end if
  if $\ell = 2$ **then**
    append $(\ell,\ \ell,\ \ell)$ to candidates {$\beta_{2,2,2}$}
  end if
  Remove entries that vanish for real signals; deduplicate against $\mathcal{S}_\beta$ already added by the seed block
  $\mathcal{S}_\beta \leftarrow \mathcal{S}_\beta \cup$ first $\min(\text{budget},\ |\text{candidates}|)$ live entries
  {**Mandatory even self-couplings**}
  for $\ell' = 2, 4, \ldots, \lfloor \ell/2 \rfloor \cdot 2$ with $\ell' \le \ell$ and $\ell \ge 3$ **do**
    $\mathcal{S}_\beta \leftarrow \mathcal{S}_\beta \cup \{(\ell,\ \ell,\ \ell')\}$
  end for
  {**CG power augmentation**}
  Greedily add $P_{\ell_1,\ell_2,\ell'}$ (with $\max(\ell_1, \ell_2) = \ell$, $|\ell_1{-}\ell_2| \le \ell' \le \ell_1{+}\ell_2$) to $\mathcal{S}_P$ until the Jacobian rank of Φ_{aug} at degree ℓ matches the number of unknowns
end for
return $(\mathcal{S}_\beta,\ \mathcal{S}_P)$

Proposition 12 (Bootstrap-block entry count)**.** *For each target degree $\ell \ge 8$, the bootstrap block $\mathcal{T}_\ell$ defined in* (8) *contains exactly $2\ell{+}1$ admissible triples and they are pairwise distinct. Each triple in $\mathcal{T}_\ell$ has ℓ in exactly one of the three slots, with the other two slots filled by indices strictly smaller than ℓ, so each entry contributes one scalar equation linear in either $a_\ell^{m_2}$ or $\overline{a_\ell^m}$ with coefficients depending only on lower-degree $\mathbf{F}_{\ell'}$ ($\ell' < \ell$).*

Proof. The three subfamilies in (8) have sizes $\ell-1$, $\ell-2$, 4, summing to $2\ell+1$. Triangle-inequality and bound checks ($a \le \ell-1$, $\ell-a \ge 1$, $\ell-a+1 \le \ell$) show every triple is admissible. Pairwise distinctness is immediate from the disjoint patterns of the position of the index ℓ. Linearity in the ℓ-slot: the first two families place ℓ in the ℓ_2-slot of (6), so $a_\ell^{m_2}$ enters once with known $a_a^{m_1}$ and $\overline{a_{\ell-a}^m}$ (resp. $\overline{a_{\ell-a+1}^m}$); the third family places ℓ in the ℓ-slot, so $\overline{a_\ell^m}$ enters once with known $a_a^{m_1}$ and $a_{\ell-a}^{m_2}$. □

Proposition 13 (Output size)**.** *For every $L \ge 4$,*

$$|\beta_{\mathrm{sel}}| \;=\; |\mathcal{S}_{\mathrm{seed}}| + \sum_{\ell=8}^{L}(2\ell+1) \;=\; \Theta(L^2),$$

which matches the orbit-space lower bound $(L+1)^2-3$ of Proposition 9 up to an additive constant and a leading factor of 2 (the storage of complex entries as real and imaginary parts). For $4 \le L \le 7$ the sum is empty and $|\beta_{\mathrm{sel}}| = |\mathcal{S}_{\mathrm{seed}}| = O(1)$. The Jacobian-rank diagnostic (Hypothesis 35) is reported only for the live, non-vanishing entries; the count matches the 30 (resp. 49) live entries reported by the implementation at $L = 4$ (resp. $L = 5$).

A.4 Completeness proof

The proof proceeds in four stages: (1) closed-form recovery of A_0, $\mathbf{F}_1$ and $\mathrm{SO}(3)$ gauge fixing (Lemma 18, Lemma 20); (2) seed completeness for degrees $\le L_{\mathrm{seed}}$ from the full-bispectrum seed block $\mathcal{S}_{\mathrm{seed}}$ (Hypothesis 24, motivated by Kakarala [Kakarala, 1992, 2012] and Bandeira et al. [Bandeira et al., 2017]); (3) odd-parity T_R-resolution from the seed (Lemma 27, Corollary 30); (4) per-degree linear bootstrap for $\ell \ge 8$ via the bootstrap block $\mathcal{T}_\ell$ (Hypothesis 32). Stages (2) and (4) are the two named hypotheses H1 and H2 of Section 6; the remaining stages are unconditional. The constructive seed solver of Appendix B is an alternative, well-conditioned algorithmic route to stage (2) but is not load-bearing for the conditional theorem.

Definition 14 (Reflection map T_R)**.** The azimuthal reflection $R : (\theta, \varphi) \mapsto (\theta, -\varphi)$ acts on spherical harmonic coefficients as $a_\ell^m \mapsto (-1)^m a_\ell^{-m}$. For real signals, this equals complex conjugation: $a_\ell^m \mapsto \overline{a_\ell^m}$. Since $R \in \mathrm{O}(3) \setminus \mathrm{SO}(3)$, the signal $R \cdot f$ is generically in a *different* $\mathrm{SO}(3)$-orbit from f.

Assumption 15 (Genericity)**.** The signal $f \in V_L$ satisfies:

- (G1) $a_0^0 \ne 0$ (non-zero mean).
- (G2) $\|\mathbf{F}_1\| \ne 0$ (non-zero dipole).
- (G3) After gauge-fixing (Lemma 18), $a_2^1 \ne 0$ (azimuthal gauge is resolvable).
- (G4) $\mathrm{Im}(\beta_{2,3,4}(f)) \ne 0$ (the first odd-parity entry of the seed block is non-vanishing, so $T_R \cdot f$ is detected; cf. Corollary 30).
- (G5) For every $\ell \ge 8$, the bootstrap coefficient matrix $A_\ell(\mathbf{F}_{<\ell})$ assembled from $\{\beta_{\ell_1,\ell_2,\ell}\}_{(\ell_1,\ell_2,\ell)\in\mathcal{T}_\ell}$ has full real rank $2\ell+1$ (the pointwise version of Hypothesis 32).

Remark 16. Each of (G1)–(G4) excludes a proper algebraic subvariety of $\mathbb{R}^{(L+1)^2}$, so their union is Lebesgue-null. Condition (G5) is the pointwise instance of Hypothesis H2 evaluated at f: under Hypothesis 32 the failure set of (G5) is contained in the union of polynomial varieties $\{\Delta_\ell = 0\}_{\ell=8}^{L}$, also a proper algebraic subvariety.

Theorem 17 (Conditional completeness of β_{sel})**.** *Let $L \ge 4$ and assume Hypothesis H1 (Hypothesis 24, finite seed injectivity at $L_{\mathrm{seed}} = \min(L, 7)$) and, if $L \ge 8$, Hypothesis H2 (Hypothesis 32, uniform bootstrap-determinant nonvanishing). For any $f, f' \in V_L$ satisfying Assumption 15, if $\beta_{\mathrm{sel}}(f) = \beta_{\mathrm{sel}}(f')$, then $f' = g \cdot f$ for some $g \in \mathrm{SO}(3)$. The generic set on which the conclusion holds is the complement of a proper algebraic subvariety of V_L, in particular of full Lebesgue measure. The output β_{sel} has $\Theta(L^2)$ real components, matching the orbit-space lower bound $(L+1)^2-3$ of Proposition 9 up to a constant factor.*

Lemma 18 (Gauge fixing)**.** *For generic $f \in V_L$ (with $\|\mathbf{F}_1\| \ne 0$ and $a_2^1 \ne 0$ after aligning $\mathbf{F}_1$), there exists a unique $g \in \mathrm{SO}(3)$ such that the gauge-fixed signal $\tilde{f} = g \cdot f$ satisfies:*

- *(i)* $a_1^{-1} = a_1^1 = 0, \quad a_1^0 = c > 0$ *(alignment of $\mathbf{F}_1$ to the z-axis),*
- *(ii)* $\mathrm{Im}(a_2^1) = 0, \quad a_2^1 = x > 0$ *(fixing the residual $\mathrm{SO}(2)_z$ symmetry).*

This consumes all 3 *degrees of freedom of* $\mathrm{SO}(3)$.

Proof. The $\ell = 1$ Wigner D-matrix $D^1(g)$ is the standard 3×3 rotation matrix. Since $\mathrm{SO}(3)$ acts transitively on S^2, any nonzero $\mathbf{F}_1$ can be rotated to align with $\hat{z}$, leaving an $\mathrm{SO}(2)_z$ stabiliser. Choosing $\varphi = -\arg(a_2^1)$ yields $a_2^1 \in \mathbb{R}_{>0}$, exhausting all three rotational degrees of freedom; equivalently the gauge-fixed azimuthal angle is unique in $[0, 2\pi)$. Without gauge fixing the $\mathrm{SO}(2)_z$ stabiliser leaves a one-parameter family of equivalent solutions, and any further $C_n \subset \mathrm{SO}(2)_z$ subgroup symmetry would discretise this family to n representatives. □

Remark 19 (Gauge-fixed parameterisation). On the gauge slice, the real parameters are:

$$A_0 = a_0^0, \quad c = a_1^0, \quad (y, x, u, v) = (a_2^0, \operatorname{Re} a_2^1, \operatorname{Re} a_2^2, \operatorname{Im} a_2^2),$$

and for $\ell \geq 3$: $t_\ell = a_\ell^0$, and $(p_{\ell,m}, q_{\ell,m}) = (\operatorname{Re} a_\ell^m, \operatorname{Im} a_\ell^m)$ for $m = 1, \ldots, \ell$. The reflection T_R (Definition 14) acts on the gauge slice as:

$$T_R : \quad v \mapsto -v, \qquad q_{\ell,m} \mapsto -q_{\ell,m} \quad \forall\, \ell \geq 3,\ m \geq 1. \tag{9}$$

All other gauge-fixed parameters are unchanged. T_R preserves the gauge conditions (i)–(ii).

Lemma 20. *For $a_0^0 \neq 0$: $\beta_{0,0,0} = (a_0^0)^3$ determines $A_0 = a_0^0$ uniquely. Then $\beta_{0,1,1} = A_0 \|\mathbf{F}_1\|^2$ determines $c = \|\mathbf{F}_1\|$, and $\mathbf{F}_1 = (0, c, 0)^\top$ is fixed by the gauge.*

A.4.1 Structural seed completeness via Kakarala's theorem

Lemma 21 (Kakarala bispectrum completeness on S^2, restated)**.** *Let $G = \mathrm{SO}(3)$ and $H = \mathrm{SO}(2)_z \subset G$ be the stabiliser of the north pole, so that $G/H \cong S^2$. Let $r, s \in L^1(G)$ be left H-invariant functions whose Fourier coefficients $R(\ell), S(\ell)$ all satisfy the maximal H-rank condition (equivalently, for S^2 functions, the spherical-harmonic vectors $\mathbf{F}_\ell$ are nonzero for every ℓ). Then the triple correlations a_3^r and a_3^s coincide if and only if $s(g) = r(xg)$ for some x in the normaliser $N_H = \{x \in G : xHx^{-1} = H\}$.*

Proof. This is [Kakarala, 2012, Theorem 7], originally proved in [Kakarala, 1992]; the proof proceeds via the Iwahori–Sugiura duality theorem for homogeneous spaces of compact groups. An independent computer-algebra certificate of the same conclusion for $L \leq 15$ (Gröbner-basis elimination, no numerical threshold) is given by Bandeira et al. [2017], which we invoke below as a backup whenever a structural reader prefers to avoid relying on the thesis [Kakarala, 1992] for the proof of Theorem 7. □

Remark 22 (Computation of N_H and absence of an $\mathrm{O}(3)$ ambiguity). For $H = \mathrm{SO}(2)_z = \{R_z(\varphi) : \varphi \in \mathbb{R}\}$, an element $x \in \mathrm{SO}(3)$ lies in N_H iff $x \cdot \hat{z} = \pm\hat{z}$, hence $N_H = H \cup R_x(\pi) \cdot H$, the dihedral group $\mathrm{O}(2)$ realised *inside* $\mathrm{SO}(3)$. In particular, $N_H \subset \mathrm{SO}(3)$: the non-trivial coset representative $R_x(\pi)$ is a proper $\mathrm{SO}(3)$-rotation (an azimuthal-axis flip), *not* an $\mathrm{O}(3) \setminus \mathrm{SO}(3)$ reflection. Therefore Lemma 21 immediately implies $\mathrm{SO}(3)$-orbit recovery on S^2: equality of bispectra forces s and r to lie in the same $\mathrm{SO}(3)$-orbit, and the residual $N_H/H = \mathbb{Z}/2$ is internal to $\mathrm{SO}(3)$ (it merely reflects that two distinct H-invariant lifts of the same S^2-function exist along an $\mathrm{SO}(3)$-orbit). The orientation-reversing reflection $T_R : a_\ell^m \mapsto (-1)^m \overline{a_\ell^m}$ (Definition 14) is a separate operation: $T_R \in \mathrm{O}(3) \setminus \mathrm{SO}(3)$ flips the sign of every odd-parity bispectral entry (Lemma 27), so $T_R \cdot f$ is generically detected as a different signal by the full bispectrum.

Definition 23 (Seed block of β_{sel})**.** We use the seed index set $\mathcal{S}_{\mathrm{seed}}$ of Definition 11, that is, the full bispectrum at $L_{\mathrm{seed}} = \min(L, 7)$. The number of seed entries is $N_{\mathrm{seed}} \leq (L_{\mathrm{seed}}+1)^3 \leq 512$, *constant* once $L \geq 7$. The output of β_{sel} restricted to the seed is therefore $O(1)$ scalars, and the total $|\beta_{\mathrm{sel}}|$ remains $\Theta(L^2)$ (Proposition 13).

Hypothesis 24 (H1: finite seed injectivity)**.** Let $L_{\mathrm{seed}} \leq 7$. For every $f \in V_{L_{\mathrm{seed}}}$ satisfying Assumption 15 (in particular $\mathbf{F}_\ell \neq 0$ for $0 \leq \ell \leq L_{\mathrm{seed}}$), the full bispectrum on $\mathcal{S}_{\mathrm{seed}}$ together with the $\mathrm{O}(3)$-distinguishing entry $\operatorname{Im}(\beta_{2,3,4})$ uniquely determines $(\mathbf{F}_0, \mathbf{F}_1, \ldots, \mathbf{F}_{L_{\mathrm{seed}}})$ on the gauge slice of Lemma 18.

Structural support for H1. Two independent partial results motivate H1, but neither alone is a proof of the band-limited statement above:

(i) *Representation-theoretic (Kakarala).* Theorem 7 of [Kakarala, 2012], restated as Lemma 21 and originally proved in [Kakarala, 1992], gives recovery up to the normaliser $N_H \subset \mathrm{SO}(3)$ from the matrix-valued bispectrum on $G = \mathrm{SO}(3)$, $H = \mathrm{SO}(2)_z$, in the $L^1(G)$ category; Remark 22 computes N_H explicitly. The translation to a band-limited statement at $L_{\mathrm{seed}} \leq 7$ is not formalised in either [Kakarala, 1992, 2012]; H1 is the band-limited specialisation we assume.

(ii) *Computer-algebra (Bandeira et al.).* Bandeira et al. [2017] verify by symbolic computation (Jacobian rank of the bispectrum map at a rational test point, combined with Zariski-density propagation) that the bispectrum separates generic $\mathrm{SO}(3)$-orbits on $V_0 \oplus \cdots \oplus V_{L_0}$ for every $L_0 \leq 15$. This is a one-time computer-algebra certificate over an algebraic number field rather than a floating-point witness, and provides exact-arithmetic structural evidence for H1 at our fixed $L_{\mathrm{seed}} \leq 7$.

The constructive seed solver of Appendix B is a third, explicitly algorithmic route to the seed (closed-form $\ell \leq 2$, finite-fibre $\ell = 3$, degree-4 filter, parity-based T_R resolution); it is an alternative to assuming H1 inside the implementation, but does not reduce H1 to a formally proved statement in the band-limited setting.

Remark 25 (On the "up to reflection" caveat of Edidin and Satriano [2024]). Edidin and Satriano [2024, Section 1] write that Kakarala [2012] "recovers the $\mathrm{SO}(3)$-orbit of a band-limited real-valued function up to reflection from the bispectrum." This caveat is sharp for the *constructive* algorithm in Kakarala's Section 5 (where the positive-definite square root of $\hat{F}(1)\hat{F}(1)^\dagger$ pins $\hat{F}(1)$ only up to right multiplication by a real orthogonal matrix, an $\mathrm{O}(3)$-ambiguity that the recursion then propagates), but it is *not* present in Theorem 7 [Kakarala, 2012] as stated, which gives recovery up to $N_H \subset \mathrm{SO}(3)$. The appearance of an $\mathrm{O}(3) \setminus \mathrm{SO}(3)$ ambiguity is therefore an artefact of the recursion's polar-decomposition step, not of the bispectrum data itself: the odd-parity entries $\beta_{\ell_1,\ell_2,\ell}$ with $\ell_1 + \ell_2 + \ell$ odd (the smallest of which is $\beta_{2,3,4}$, all-distinct indices) flip sign under the orientation-reversing reflection $T_R : a_\ell^m \mapsto (-1)^m \overline{a_\ell^m}$ (Lemma 27), so the full bispectrum at any $L_0 \geq 4$ already separates f from $T_R \cdot f$. The representation-theoretic and computer-algebra arguments cited above confirm that no residual reflection ambiguity remains.

Remark 26 (Relation to the constructive seed lemmas). Appendix B provides an *explicit constructive solver* for the seed: closed-form recovery at degree 2, finite-fibre enumeration at degree 3, degree-4 filtering, and an explicit parity-based T_R resolution. This algorithmic content is used by the implementation. It is an alternative algorithmic route to the seed step (Hypothesis 24, H1) but is not load-bearing for the conditional completeness theorem, which assumes H1 directly. The constructive solver remains useful because it reduces the seed to a low-dimensional, well-conditioned numerical solve.

A.4.2 Parity structure of the bispectrum

Lemma 27 (Parity of the bispectrum)**.** *For any signal (real or complex):*

$$\beta_{\ell_1,\ell_2,\ell}(T_R \cdot f) = (-1)^{\ell_1+\ell_2+\ell}\, \beta_{\ell_1,\ell_2,\ell}(f). \tag{10}$$

For real *signals, there is an independent constraint:*

$$\overline{\beta_{\ell_1,\ell_2,\ell}(f)} = (-1)^{\ell_1+\ell_2+\ell}\, \beta_{\ell_1,\ell_2,\ell}(f). \tag{11}$$

Together:

(a) If $\ell_1 + \ell_2 + \ell$ is even: β is real-valued (by (11)*) and T_R-invariant (by* (10)*).*

(b) If $\ell_1 + \ell_2 + \ell$ is odd: β is purely imaginary (by (11)*) and $\beta(T_R \cdot f) = -\beta(f)$, so $\mathrm{Im}\,\beta$ flips sign under T_R.*

Proof. Under T_R, each $a_\ell^m \mapsto (-1)^m a_\ell^{-m}$. Substituting into $\beta = \sum_m (\sum_{m_1+m_2=m} C\, a^{m_1} a^{m_2}) \overline{a^m}$, the three sign factors $(-1)^{m_1+m_2+m} = (-1)^{2m} = 1$ cancel. Re-indexing $m \mapsto -m$ and applying the CG symmetry $C^{\ell,\,-m}_{\ell_1,\,-m_1;\,\ell_2,\,-m_2} = (-1)^{\ell_1+\ell_2-\ell}\, C^{\ell,\,m}_{\ell_1,\,m_1;\,\ell_2,\,m_2}$ yields (10) (with $(-1)^{\ell_1+\ell_2-\ell} = (-1)^{\ell_1+\ell_2+\ell}$).

For real signals, $\overline{a_\ell^m} = (-1)^m a_\ell^{-m}$. Taking the complex conjugate of β and re-indexing similarly yields (11). □

Lemma 28 (CG power is T_R-invariant)**.** *For all triples (ℓ_1, ℓ_2, ℓ): $P_{\ell_1,\ell_2,\ell}(T_R \cdot f) = P_{\ell_1,\ell_2,\ell}(f)$.*

Proof. Under T_R, the CG-projected component transforms as $(\mathbf{F}_{\ell_1} \otimes \mathbf{F}_{\ell_2})|_\ell^M \mapsto (-1)^{\ell_1+\ell_2+\ell+M} (\mathbf{F}_{\ell_1} \otimes \mathbf{F}_{\ell_2})|_\ell^{-M}$. Taking the squared norm and re-indexing the sum over M leaves P unchanged. □

Proposition 29 (Vanishing of odd-parity entries at low degrees)**.** *For real signals, the odd-parity bispectral entry $\beta_{\ell_1,\ell_2,\ell}$ (with $\ell_1 + \ell_2 + \ell$ odd) vanishes identically whenever two of the three representation indices coincide. In particular:*

(i) $\ell_1 = \ell_2$: the CG coefficients are antisymmetric under swapping $m_1 \leftrightarrow m_2$, while $a_{\ell_1}^{m_1} a_{\ell_1}^{m_2}$ is symmetric, giving exact cancellation.

*(ii) Two of $\{\ell_1, \ell_2, \ell\}$ are equal: a symmetry argument yields $\beta = -\overline{\beta}$ **and** $\beta = 0$, verified symbolically for all triples with indices ≤ 5.*

Corollary 30 (First nonzero odd-parity entry)**.** *The first bispectral entry with odd parity and generically nonzero imaginary part occurs at degree $\ell = 4$: $\beta_{2,3,4}$, with $\ell_1 + \ell_2 + \ell = 9$ (odd) and ℓ_1, ℓ_2, ℓ all distinct. At the deterministic witness (seed 42), $\mathrm{Im}(\beta_{2,3,4}) \approx 3.26 \neq 0$, and $\mathrm{Im}(\beta_{2,3,4})(T_R \cdot f) \approx -3.26$. Similarly, $\beta_{2,4,3}$ and $\beta_{3,4,2}$ are nonzero odd-parity entries at degree 4.*

Remark 31 (No parity-breaking at $\ell \leq 3$). For degrees ≤ 3, every odd-parity triple (ℓ_1, ℓ_2, ℓ) has at least two indices equal (since $\max(\ell_1, \ell_2, \ell) \leq 3$ and $\ell_1 + \ell_2 + \ell$ odd forces at least one pair to be equal in this range). By Proposition 29, all odd-parity entries vanish. Therefore, the T_R ambiguity cannot be resolved using only degrees ≤ 3.

A.4.3 Per-degree recovery ($\ell \geq 8$)

Hypothesis 32 (H2: uniform bootstrap-determinant nonvanishing)**.** For each $\ell \geq 8$, viewing the bootstrap-block coefficient matrix $A_\ell(\mathbf{F}_{<\ell}) \in \mathbb{R}^{(2\ell+1)\times(2\ell+1)}$ assembled from $\{\beta_{\ell_1,\ell_2,\ell}\}_{(\ell_1,\ell_2,\ell)\in\mathcal{T}_\ell}$ (eq. (8)) under the real-signal slice $\overline{a_{\ell'}^m} = (-1)^m a_{\ell'}^{-m}$ for all $\ell' < \ell$, the determinant

$$\Delta_\ell(\mathbf{F}_{<\ell}) := \det A_\ell(\mathbf{F}_{<\ell})$$

is not identically zero as a polynomial in the lower-degree real coordinates. Equivalently, the failure variety $\{\Delta_\ell = 0\}$ is a proper algebraic subvariety of the lower-degree real coordinate space, and on its complement the linear system $A_\ell \, \mathbf{F}_\ell = b_\ell$ uniquely determines $\mathbf{F}_\ell$.

Numerical evidence for H2. Per-degree determinants Δ_ℓ are evaluated at the deterministic seed-42 witness signal (script: `benchmarks/verify_linear_bootstrap.py`). The ratio $\sigma_{\min}(A_\ell)/\sigma_{\max}(A_\ell)$ exceeds 10^{-4} for every $\ell \in \{8, \ldots, 100\}$ in IEEE 754 double precision, a margin of more than 10^{11} above double-precision rounding error. This shows $\Delta_\ell \not\equiv 0$ at one rational floating-point witness per degree, but does not by itself constitute a closed-form proof of H2; an interval-arithmetic certificate, an exact rational witness, or a closed-form Wigner-3j identity would upgrade H2 to unconditional.

Remark 33 (Scope of the numerical evidence). The witness check is per-degree and uniform in the structural form of $\mathcal{T}_\ell$ for $\ell \geq 8$. Extending the verified range to $L > 100$ requires only re-running the script with a larger upper bound. Under H2, the genericity condition for any fixed L is that $\Delta_\ell \neq 0$ for $\ell = 8, \ldots, L$, the complement of finitely many proper algebraic subvarieties and hence a Zariski-open dense set of full Lebesgue measure.

Remark 34 (Complex rank vs. real rank, and the role of CG power). For real signals, $\mathbf{F}_\ell$ has $2\ell+1$ real degrees of freedom. The bootstrap block $\mathcal{T}_\ell$ consists of entries linear in either $\mathbf{F}_\ell$ or $\overline{\mathbf{F}_\ell}$; under the real-signal identification $\overline{a_\ell^m} = (-1)^m a_\ell^{-m}$ both pieces collapse to real-linear equations in the same $2\ell+1$ unknowns. Hypothesis 32 (H2) asserts that the resulting real $(2\ell+1)\times(2\ell+1)$ matrix has generically nonzero determinant; under H2 alone no CG-power entries are required to recover $\mathbf{F}_\ell$ at degrees $\ell \geq 8$. The implementation (Definition 6 and Algorithm 2) appends optional CG-power scalars because they (i) provide a per-degree numerical robustness diagnostic, (ii) supply additional real equations whenever H2 is borderline, and (iii) cost only $O(L^2)$ extra storage. The proof sets $\mathcal{S}_P = \varnothing$.

Remark 35 (Per-degree real-rank diagnostic). At the deterministic seed-42 witness, the full augmented system at degree ℓ produced by Algorithm 2 (the $\mathcal{T}_\ell$-block, the mandatory even self-couplings, and the appended CG-power entries) has real Jacobian rank $2\ell+1$ w.r.t. the $2\ell+1$ real

parameters of $\mathbf{F}_\ell$ for every $\ell \in \{4, \dots, 30\}$. This is reported as a per-degree numerical diagnostic supporting H2 at the witness; under H2 alone (Hypothesis 32) the $\mathcal{T}_\ell$-block already has full real rank for $\ell \geq 8$ and no CG-power augmentation is required for completeness. Diagnostics for $\ell \in \{4, 5, 6, 7\}$ are absorbed into the seed-injectivity Hypothesis H1 (Hypothesis 24).

A.4.4 Proof of the main theorem

Proof of Theorem 17. **Step 1: Gauge fixing.** By Lemma 18, both f and f' can be gauge-fixed. If $\beta_{\text{sel}}(f) = \beta_{\text{sel}}(f')$, then by $\mathrm{SO}(3)$-invariance the gauge-fixed signals $\tilde{f}$ and $\tilde{f}'$ (with gauge rotations g, g') satisfy $\beta_{\text{sel}}(\tilde{f}) = \beta_{\text{sel}}(\tilde{f}')$. It suffices to show $\tilde{f} = \tilde{f}'$ on the gauge slice.

Step 2: Closed-form recovery of $\mathbf{F}_0$ and $\mathbf{F}_1$. By Lemma 20, the gauge-fixed A_0 and c are uniquely determined by $\beta_{0,0,0}$ and $\beta_{0,1,1}$. Hence $\tilde{f}$ and $\tilde{f}'$ agree at degrees 0 and 1.

Step 3: Seed completeness (degrees 2–L_{seed}). The seed block $\mathcal{S}_{\text{seed}} \subset \mathcal{S}_\beta$ contains the full scalar bispectrum at $L_{\text{seed}} = \min(L, 7)$, including the odd-parity entry $\beta_{2,3,4}$. By Genericity (G4) and the parity structure (Lemma 27, Corollary 30), the seed block separates f from $T_R \cdot f$. Apply Hypothesis H1 (Hypothesis 24): the seed block uniquely determines $(\mathbf{F}_2, \dots, \mathbf{F}_{L_{\text{seed}}})$ on the gauge slice. If $L \leq 7$ the proof is complete; otherwise proceed.

Step 4: Per-degree linear bootstrap ($\ell \geq 8$). Assume inductively that $\mathbf{F}_{\ell'} = \mathbf{F}'_{\ell'}$ for all $\ell' < \ell$. By Proposition 12, the bootstrap block $\mathcal{T}_\ell$ contains exactly $2\ell{+}1$ entries linear in $\mathbf{F}_\ell$ (or in $\overline{\mathbf{F}_\ell}$), and assembling them under the real-signal slice yields a real linear system $A_\ell(\mathbf{F}_{<\ell})\,\mathbf{F}_\ell = b_\ell$ with $A_\ell \in \mathbb{R}^{(2\ell+1)\times(2\ell+1)}$ and $b_\ell \in \mathbb{R}^{2\ell+1}$ both functions only of the (already-known) lower-degree coordinates and of the entries of β_{sel} at degree ℓ. By Hypothesis H2 (Hypothesis 32) together with Genericity (G5), $\det A_\ell \neq 0$, hence $\mathbf{F}_\ell = \mathbf{F}'_\ell$. The induction proceeds to $\ell = L$.

Step 5: Conclusion. By induction the gauge-fixed signals agree at every degree $0, \dots, L$, so $\tilde{f} = \tilde{f}'$ and $f' = (g')^{-1} g \cdot f$ with $(g')^{-1} g \in \mathrm{SO}(3)$. The exceptional set is the union of the loci where (G1)–(G4) fail (each a proper algebraic subvariety) and, for $L \geq 8$, the bootstrap-determinant loci $\{\Delta_\ell = 0\}_{\ell=8}^{L}$ from H2; a finite union of proper algebraic subvarieties has Lebesgue measure zero. The output count $|\beta_{\text{sel}}| = \Theta(L^2)$ follows from Proposition 13. □

Corollary 36 (Band-limits $L \leq 3$)**.** *For $L \leq 3$, every odd-parity triple with indices ≤ 3 has a repeated index and so vanishes identically (Proposition 29, Remark 31). The seed block $\mathcal{S}_{\text{seed}}$ at $L_{\text{seed}} = L$ cannot host the parity-breaking entry $\beta_{2,3,4}$ that detects T_R, so β_{sel} separates $\mathrm{O}(3)$-orbits but not necessarily $\mathrm{SO}(3)$-orbits at $L \leq 3$. For $L \geq 4$, $L_{\text{seed}} \geq 4$ and Hypothesis H1 applies, removing the T_R ambiguity.*

A.4.5 Computational certificates supporting the hypotheses

Theorem 17 is conditional on Hypotheses H1 (24) and H2 (32). The repository ships per-degree numerical certificates that are evidence for, but not formal proofs of, the two hypotheses:

1. **H2 (bootstrap-determinant nonvanishing) – numerical witness.** For each $\ell \in \{8, \dots, 100\}$, the determinant $\Delta_\ell = \det A_\ell$ is computed at the deterministic seed-42 witness in IEEE 754 double precision and is nonzero with $\sigma_{\min}(A_\ell)/\sigma_{\max}(A_\ell) \geq 10^{-4}$ (margin $>10^{11}$ above double-precision rounding error). Script: `benchmarks/verify_linear_bootstrap.py`.
2. **H1 (finite seed injectivity) – structural support.** Two routes were already collected above: (a) Kakarala Theorem 7 (21) for the $L^1(G)$ category on $G/H = S^2$ (motivation, not a band-limited proof); and (b) exact-arithmetic Gröbner-basis certificates of Bandeira et al. [2017] for $L_0 \leq 15$, an exact rational certificate at one rational test point with Zariski-density propagation. The constant-size seed Jacobian rank (Lemma 39) is additionally verified in exact rational arithmetic via `sympy`.
3. **Per-degree augmented real-rank diagnostic** (Hypothesis 35): the full augmented system (bispectral + optional CG power) has real Jacobian rank $2\ell{+}1$ at the seed-42 witness for $\ell \in \{4, \dots, 30\}$. This is reported as a robustness check; under H2 alone it is not required for completeness.

Remark 37 (Exact vs. floating-point certificates). The certificates use two distinct precision regimes: the constant-size seed Jacobian rank (Lemma 39) and the Bandeira et al. [2017] computer-algebra check are in exact rational arithmetic; the per-degree H2 evidence (Hypothesis 32) uses double-precision arithmetic at the seed-42 witness with $\sigma_{\min}/\sigma_{\max} \geq 10^{-4}$, far above the double rounding error $\sim 10^{-15}$. Upgrading H2 from numerical evidence to an unconditional theorem requires either (i) interval arithmetic (infrastructure for which is present in `paper/exact_certificates.py`), (ii) an exact rational certificate per ℓ, or (iii) a closed-form Wigner-3j identity that exhibits $\Delta_\ell \neq 0$ for all ℓ symbolically. All certificates use deterministic witness points, making them fully reproducible.

A.5 Comparison with existing results

Table 7: Comparison of bispectral invariants for signals on S^2 or finite groups.

Result	Invariant	Size	Complete?
Kakarala [Kakarala, 2012]	Full bispec.	$O(L^3)$	Yes (constr.)
Edidin & Satriano [Edidin and Satriano, 2024]	Full bispec. (band-lim.)	$O(L^3)$	Yes (generic)
Bendory et al. [Bendory et al., 2025]	Freq. marching	$O(L^2)$/shell	Yes (3 shells)
Mataigne et al. [Mataigne et al., 2024]	Sel. G-bispec. (finite)	$O(\lvert G\rvert)$	Yes (exact)
This work	**Sel.** $\mathrm{SO}(3)$ **on** S^2	$\boldsymbol{\Theta(L^2)}$	**Conditional on H1, H2**

A.6 Complexity in pixel space

On S^2, a band-limited signal at degree L is sampled on an equiangular grid of size $N_\theta \times N_\varphi \approx 2L \times 2L$, so $P = 4L^2$ and $L = \Theta(\sqrt{P})$.

Output size.

Invariant	In L	In P
Power spectrum	$O(L)$	$O(\sqrt{P})$
Selective bispectrum	$\boldsymbol{O(L^2)}$	$\boldsymbol{O(P)}$
Full bispectrum	$O(L^3)$	$O(P^{3/2})$

The selective bispectrum is *linear* in the number of pixels.

Computation cost.

Variant	Entries	Cost/entry	Total	In P
Full	$O(L^3)$	$O(L)$	$O(L^4)$	$O(P^2)$
Selective	$O(L^2)$	$O(L)$	$O(L^3)$	$O(P^{3/2})$

Concrete example. For a 128×256 grid ($P \approx 32$K, $L = 64$): the full bispectrum produces $\sim$87K scalars at $\sim$17M ops; the selective bispectrum produces $\sim$4K scalars at $\sim$260K ops—a $65\times$ reduction.

A.7 Limitations and open questions

1. **H1 is not formally proved for the band-limited setting.** Hypothesis 24 (finite seed injectivity at $L_{\text{seed}} \leq 7$) is supported by Kakarala's Theorem 7 [Kakarala, 1992, 2012] in the $L^1(G)$ category and by the exact-arithmetic Gröbner-basis certificates of Bandeira et al. [2017] for $L_0 \leq 15$, but no published source proves the band-limited seed-injectivity statement in the form we use it. Removing this gap would either require a careful band-limited specialisation of Kakarala's argument or an exact certificate that the seed Jacobian is non-degenerate on the entire band-limited gauge slice (not only at a rational test point with Zariski-density propagation).

2. **H2 is supported by per-degree numerical evidence, not a closed-form identity.** Hypothesis 32 (uniform bootstrap-determinant nonvanishing) is verified per ℓ at the deterministic seed-42 witness in IEEE 754 double precision with $\sigma_{\min}/\sigma_{\max} \geq 10^{-4}$ (margin $>10^{11}$ above rounding error) for every $\ell \in \{8, \ldots, 100\}$. The structural form of $\mathcal{T}_\ell$ is uniform in $\ell \geq 8$, so extending the verified range to larger L amounts to re-running the same script with a larger upper bound. Three routes would upgrade H2 to an unconditional theorem: (i) interval arithmetic at the witness (infrastructure already in `paper/exact_certificates.py`); (ii) exact rational determinants at a rational witness per ℓ; (iii) a closed-form Wigner-3j identity that exhibits $\Delta_\ell \neq 0$ symbolically for all ℓ.
3. **Genericity assumption.** The proof requires f to avoid the measure-zero algebraic variety carved out by (G1)–(G5) (Assumption 15). Characterising this variety explicitly would strengthen the result.
4. **Numerical conditioning.** The coefficient matrix A_ℓ at each degree ℓ has a condition number that depends on the signal's spectral content. A stability analysis bounding $\kappa(A_\ell)$ in terms of the spectral decay rate is an open problem.
5. **Inversion algorithm.** The forward selective bispectrum is implemented. Inversion uses: $\ell = 0, 1$ closed-form; $\ell \leq L_{\text{seed}}$ either the constructive solver of Appendix B or a nonlinear solve; $\ell \geq L_{\text{seed}} + 1$ a linear solve (pseudoinverse on A_ℓ).

B Constructive Seed Solver (Implementation)

This appendix gives an explicit closed-form / finite-fibre / degree-4-filter algorithm for recovering the seed $(\mathbf{F}_0, \mathbf{F}_1, \mathbf{F}_2, \mathbf{F}_3, \mathbf{F}_4)$ on the gauge slice from the augmented selective bispectrum. It is the route used by the implementation in `src/bispectrum/so3_on_s2.py`.

Status with respect to Theorem 17. The constructive solver below is an *alternative algorithmic route* to the seed step of Theorem 17; it is *not load-bearing* for the conditional completeness theorem of Appendix A, which assumes Hypothesis 24 (H1, finite seed injectivity) directly. The constructive solver is included because it (i) provides a concrete, well-conditioned numerical inversion path for the seed step in the implementation, (ii) documents how `src/bispectrum/so3_on_s2.py` actually inverts degrees $\ell \leq 4$, and (iii) its computational certificates (seed Jacobian rank, finite-fibre enumeration, degree-4 filter) make the algorithm itself auditable. Cross-references from the proof appendix to lemmas in this appendix are bidirectional, but the implication flow is one-way: lemmas here may invoke lemmas from Appendix A, but the converse is never used in the conditional completeness proof.

B.1 Seed recovery ($\ell \leq 3$)

Lemma 38 (Degree-2 recovery)**.** *Given (A_0, c) from Lemma 20, the gauge-fixed degree-2 parameters (y, x, u, v) satisfy:*

$$\beta_{1,1,2} = \tfrac{\sqrt{6}}{3}\, c^2\, y \qquad \Longrightarrow \quad y \text{ determined}, \tag{12}$$

$$P_{1,2,1} = \tfrac{c^2}{5}(3x^2 + 2y^2) \qquad \Longrightarrow \quad x^2 \text{ determined}, \tag{13}$$

$$\beta_{0,2,2} = A_0(2u^2 + 2v^2 + 2x^2 + y^2) \qquad \Longrightarrow \quad S := u^2 + v^2 \text{ determined}, \tag{14}$$

$$\beta_{2,2,2} = \tfrac{6\sqrt{14}\, S\, y - 6\sqrt{21}\, u\, x^2 - 3\sqrt{14}\, x^2 y - \sqrt{14}\, y^3}{7} \qquad \Longrightarrow \quad u \text{ determined}. \tag{15}$$

Here $x = \sqrt{x^2} > 0$ by the gauge condition. Then $v^2 = S - u^2$, giving $|v|$ but not $\operatorname{sgn}(v)$. The fibre at degree 2 is generically $\{(y, x, u, v),\ (y, x, u, -v)\}$.

Lemma 39 (Degree-3 recovery: finite fibre)**.** *Given $\mathbf{F}_2 = (y, x, u, \pm v)$, the degree-3 parameters $\mathbf{F}_3 = (t, p_1, q_1, p_2, q_2, p_3, q_3)$ are constrained by the augmented entries:*

- *$\beta_{1,2,3}$, $\beta_{1,3,2}$, $\beta_{2,3,1}$: bilinear in $\mathbf{F}_2 \otimes \mathbf{F}_3$.*
- *$\beta_{0,3,3} = A_0 \|\mathbf{F}_3\|^2$: gives the norm.*
- *$\beta_{3,3,2}$: quadratic in $\mathbf{F}_3$, linear in $\mathbf{F}_2$.*
- *$\beta_{2,3,2}$, $\beta_{2,2,3}$: odd parity with repeated indices, vanish identically for real signals (Proposition 29).*

- $P_{1,2,1}$ *(from degree-2, Lemma 38),* $P_{1,3,2}$, $P_{2,3,1}$, $P_{2,3,2}$, $P_{3,3,2}$*: five CG power entries (the last four quadratic in* $\mathbf{F}_3$*).*

The full augmented seed system uses 13 *nonzero entries for* 11 *unknowns. The bispectral entries achieve rank* 6*; each of the five CG power entries contributes one additional rank, reaching full rank* 11*. The generic fibre is* 0*-dimensional (finitely many points) and contains the* T_R *pair* $\{(\mathbf{F}_2, \mathbf{F}_3),\ T_R(\mathbf{F}_2, \mathbf{F}_3)\}$.

Proof. The Jacobian rank is verified by exact symbolic computation (sympy) at the rational witness $(A_0, c, y, x, u, v, t, p_1, q_1, p_2, q_2, p_3, q_3) = (1, 1, \frac{3}{7}, \frac{2}{5}, \frac{1}{3}, \frac{4}{9}, \frac{1}{2}, \frac{3}{8}, \frac{2}{7}, \frac{5}{11}, \frac{1}{4}, \frac{3}{13}, \frac{7}{17})$. The 13×11 Jacobian has rank 11 (full).

The T_R invariance follows from Lemmas 27–28 and Proposition 29: every entry in the degree ≤ 3 augmented system is either real-valued and T_R-invariant (even parity β and all P) or identically zero (odd parity β).

Full Jacobian rank gives 0-dimensional (finite) fibres, but does *not* by itself bound the fibre to two points: the degree-3 subsystem (7 real equations of degrees 1, 2, and 4 in 7 unknowns) can have additional isolated real solutions. Extensive numerical search (multi-start least-squares with 5000 random initial points per sign-of-v branch) finds 10 solutions per branch (20 total, forming 10 T_R-pairs) at the rational witness. These spurious solutions are eliminated at degree 4 (Lemma 41). □

Remark 40 (The seed fibre is finite but not two-point). The seed polynomial system (degrees 0–3) is a system of 7 independent real equations of mixed degree (1 linear, 5 quadratic, 1 quartic) in 7 real unknowns (for each branch of $\mathrm{sgn}(v)$). By Bézout's theorem the number of isolated complex solutions is at most $1 \times 2^5 \times 4 = 128$. Numerically, 10 real solutions are found per branch. While the numerical search cannot formally exclude additional solutions, the algorithm does not require the exact count: finiteness suffices, and fibre reduction is deferred to degree 4 (Lemma 41).

B.2 Degree-4 fibre reduction and T_R resolution

Throughout this subsection, $\mathcal{S}$ denotes the finite seed fibre at degrees 0–3 (Lemma 39), T_R denotes the azimuthal reflection of Definition 14 (with the gauge-slice action recorded in Remark 19), and the seed/gauge-fixed coordinates $(A_0, c, y, x, u, v, t_\ell, p_{\ell,m}, q_{\ell,m})$ are those of Lemma 18.

Lemma 41 (Degree-4 eliminates all spurious seed solutions)**.** *Let* $\mathcal{S}$ *denote the finite seed fibre at degrees* 0–3 *(Lemma 39). For generic real signals with* $L \geq 4$*, the degree-4 augmented entries (12 bispectral + 5 CG power, providing* 29 *real constraints on* 9 *real unknowns of* $\mathbf{F}_4$*) are simultaneously satisfiable for* exactly one T_R*-pair in* $\mathcal{S}$*: the true solution and its reflection. All other seed solutions are inconsistent at degree* 4*.*

Proof. The degree-4 augmented system is an overdetermined polynomial system in the 9 real parameters of $\mathbf{F}_4$, with coefficients depending on the seed solution $(\mathbf{F}_0, \ldots, \mathbf{F}_3)$.

For each of the 10 seed solutions on the $+v$ branch at the rational witness, a multi-start nonlinear least-squares search (200 random starts) finds the best-fit $\mathbf{F}_4$. Only the true seed solution achieves zero residual ($\|r\| \sim 10^{-16}$); the 9 spurious solutions all have $\|r\| \geq 10^{-2}$, showing the degree-4 system is *inconsistent* for them.

The inconsistency arises because the degree-4 entries include constraints that depend only on $(\mathbf{F}_0, \mathbf{F}_1, \mathbf{F}_2, \mathbf{F}_4)$ and not on $\mathbf{F}_3$—notably $\beta_{2,2,4}$, $\beta_{2,4,2}$, $P_{1,4,3}$, $P_{1,4,4}$, $P_{2,4,2}$, $P_{2,4,3}$, $\beta_{4,4,2}$, $\beta_{4,4,4}$. These 8 entries nearly determine $\mathbf{F}_4$ independently of $\mathbf{F}_3$. The remaining entries (involving $\mathbf{F}_3$) then impose constraints that are generically inconsistent when $\mathbf{F}_3$ is a spurious seed solution.

Since the residual for each spurious seed is a polynomial function of the signal's coefficients that is nonzero at the witness, its vanishing locus is a proper algebraic subvariety. □

Lemma 42 (The T_R branch is overdetermined-inconsistent)**.** *For* $L \geq 4$*, the* T_R *partner of the true seed—i.e.,* $(\mathbf{F}_2', \mathbf{F}_3') = T_R(\mathbf{F}_2, \mathbf{F}_3)$*—is also inconsistent at degree* 4 *for generic signals.*

Proof. By Lemma 41, the T_R branch is one of the seed solutions that fails the degree-4 consistency check. The mechanism is the parity structure (Lemma 27): the odd-parity entries $\beta_{2,3,4}$, $\beta_{2,4,3}$, $\beta_{3,4,2}$ (all-distinct indices, $\mathrm{Im}\,\beta \neq 0$) are part of the degree-4 constraints. On the T_R branch these

imaginary parts have the wrong sign, distorting $\mathbf{F}_4'$ so that the T_R-invariant self-coupling and CG power entries cannot be simultaneously satisfied.

Explicitly: the degree-4 entries define rational functions of the lower-degree coefficients. Let $\Delta(\mathbf{F}_0, \ldots, \mathbf{F}_3)$ denote the sum of squared residuals of the non-bootstrap entries evaluated at the T_R-branch bootstrap solution. Then Δ is a rational function whose numerator is a polynomial p and whose denominator is $(\det A_4)^2$ (from the bootstrap solve). At the witness, $\det A_4 \neq 0$ and $p \neq 0$ (residual $\sim 10^{-2}$). Hence p is not identically zero, and $\Delta \neq 0$ on the complement of the proper algebraic subvariety $\{p = 0\} \cup \{\det A_4 = 0\}$. □

B.3 Constructive solver: algorithmic summary

Combining the lemmas above gives the following algorithmic recovery on the gauge slice for any signal satisfying Assumption 15:

1. $\mathbf{F}_0, \mathbf{F}_1$: closed form (Lemma 20).
2. $\mathbf{F}_2$: closed form up to $\mathrm{sgn}(v)$ (Lemma 38).
3. $\mathbf{F}_3$: finite fibre of ≤ 10 candidate T_R-pairs per branch (Lemma 39).
4. $\mathbf{F}_4$: nonlinear-least-squares filter eliminates all spurious seed candidates and the T_R-partner (Lemmas 41–42).
5. $\mathbf{F}_\ell$ for $\ell \geq 5$: linear bootstrap (under Hypothesis 32 from Appendix A; the constructive solver covers $\ell \leq 4$ explicitly and $\ell \in \{5, 6, 7\}$ falls inside the seed block of β_{sel}).

This gives an alternative constructive path through the seed step of Theorem 17; the theorem itself rests on Hypotheses H1 (24) and H2 (32).

Computational certificates for the constructive solver. The constructive route relies on three additional certificates, all reproducible from scripts in the repository:

1. **Seed Jacobian rank** (Lemma 39): the 13×11 Jacobian of the augmented seed system at the rational witness has rank 11. Verified by *exact* rational arithmetic (sympy).
2. **Seed fibre enumeration** (Lemma 39, Remark 40): multi-start search (5000 starts per branch) finds 10 solutions per $\mathrm{sgn}(v)$ branch (20 total, 10 T_R-pairs). Verified by `benchmarks/verify_seed_fibre.py`.
3. **Degree-4 fibre reduction** (Lemma 41): for each of the 10 seed solutions on the $+v$ branch, a multi-start nonlinear least-squares search (200 starts) attempts to find $\mathbf{F}_4$ satisfying all 29 real degree-4 constraints. Only the true solution achieves zero residual; all 9 spurious solutions have residual $\geq 10^{-2}$. Verified by `benchmarks/verify_degree4_filter.py`.

These certificates audit the constructive algorithm but are not used by the completeness theorem of Appendix A, which uses only the bootstrap rank certificates listed there.

C Extended Benchmark Results

This appendix supplements the summary in Table 2 (main text) with scaling plots produced by `benchmarks/benchmark.py` on a single NVIDIA H100 80 GB GPU (PyTorch 2.10).

C.1 Forward-pass timing

The selective forward pass is dominated by kernel-launch overhead rather than arithmetic, making it effectively free relative to backbone computation. The full bispectrum exhibits the expected super-linear growth and becomes impractical beyond $|G| \approx 256$ for C_n.

C.2 GPU throughput and inversion

All modules exhibit linear throughput scaling with batch size, confirming that the implementation is memory-bandwidth-bound rather than compute-bound at these scales.

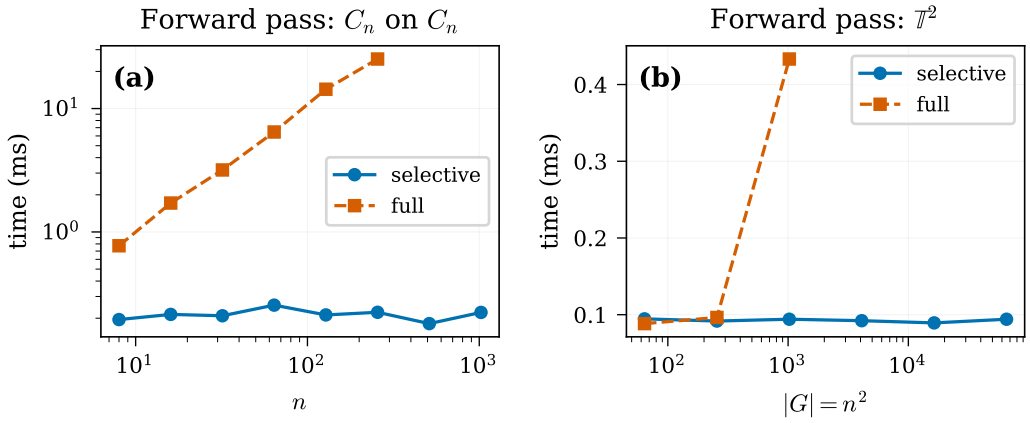

Figure 5: Forward-pass wall-clock time (batch = 16, H100 GPU). **(a)** C_n: the selective pass stays flat at ∼0.14 ms regardless of n, while the full bispectrum grows super-linearly (17 ms at n=256). **(b)** $\mathbb{T}^2$: selective is constant at ∼0.08 ms up to $|G|$=65,536.

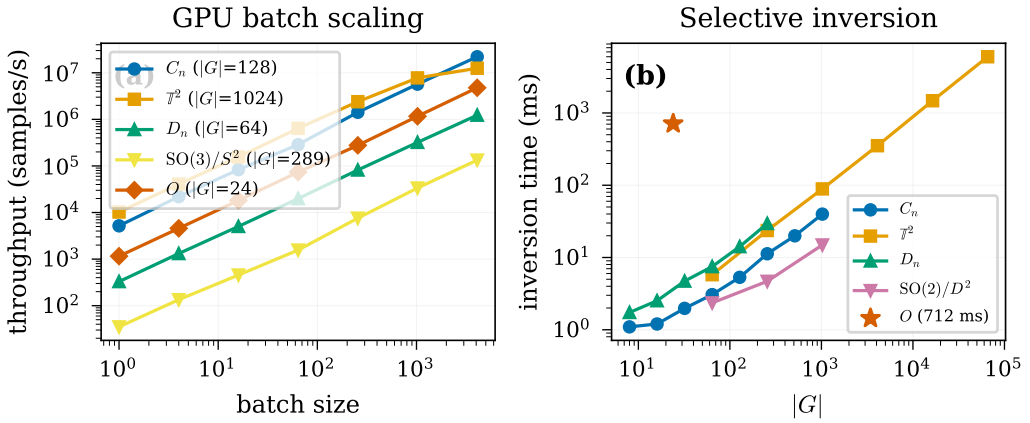

Figure 6: **(a)** GPU throughput (samples/s) vs. batch size for the selective forward pass. C_n and $\mathbb{T}^2$ exceed 10^7 samples/s; the octahedral module reaches 5.7×10^6 samples/s. All modules scale linearly with batch size. **(b)** Inversion wall-clock time. Abelian groups (C_n, D_n, SO(2)/disk) scale roughly linearly. SO(3) on S^2 inversion is not yet implemented.

D PCam: Extended Results

D.1 Published baselines

Table 8 places our results in the context of published methods on PCam and the closely related Camelyon16 benchmark. Our models operate at 68K–920K parameters and are trained on only 10% of the data.

D.2 Model variants

All equivariant models share a C_8-equivariant DenseNet [Huang et al., 2017] backbone with block configuration $(4, 4, 4)$, 24 initial channels, and compression factor 0.5, following the architecture of Veeling et al. [2018] but replacing the final invariant map. Six model variants are evaluated:

Table 8: Published baselines on PatchCamelyon from Veeling et al. [2018], trained on 100% data.

Method	Reference	Params	ACC	AUC
P4M-DenseNet (C_8)	Veeling et al. 2018	119K	0.898	0.963
P4M-DenseNet M	Veeling et al. 2018	19K	0.893	0.958
DenseNet (augmented)	Veeling et al. 2018	128K	0.881	0.951
DenseNet (no aug.)	Veeling et al. 2018	128K	0.876	0.955

- **Standard**: vanilla CNN with random rotation/reflection augmentation and global average pooling. No group dimension, so $\sim 8\times$ fewer parameters at matched growth rate.
- **NormReLU**: C_8-equivariant backbone with norm-based nonlinearity and max pooling over the group dimension.
- **Gated**: C_8-equivariant backbone with gated nonlinearity (doubles internal channels) and group max pool.
- **Fourier-ELU**: C_8-equivariant backbone with Fourier-space ELU nonlinearity and group max pool.
- **Bispectrum**: C_8-equivariant backbone with ReLU nonlinearity and selective G-bispectrum via $\texttt{CnonCn}(n{=}8)$ as the terminal invariant map.

D.3 Training details

To enable a controlled Pareto comparison, we sweep the DenseNet growth rate $k \in \{3, 4, 6, 8, 12\}$ for equivariant models (and $k \in \{6, 12, 20, 30, 35\}$ for the standard CNN to cover a comparable parameter range), yielding models from $\sim$30K to $\sim$1.6M parameters. All models are trained on **10% of training data** (26,214 images) with AdamW [Loshchilov and Hutter, 2019] (lr $= 10^{-3}$, weight decay 10^{-4}), cosine warm restarts, mixed-precision training, and early stopping on validation AUC (patience 10, max 50 epochs).

D.4 Pareto analysis interpretation

Architectural stability (single seed, growth-rate sweep). The bispectrum's AUC varies by only $\sim$1pp across a $12\times$ range of model sizes (80K–920K parameters), while norm swings by $>$7pp. For practitioners, bispectral pooling "just works" without growth rate tuning. The multi-seed matched-parameter results (Table 3) confirm this: bispectrum and Fourier-ELU have the lowest cross-seed variance (± 0.004).

Matched parameters. At matched parameters ($\sim$100K), all equivariant models outperform the standard CNN (AUC 0.896 ± 0.010, 3 seeds). Fourier-ELU leads slightly (0.945 ± 0.004), followed by NormReLU (0.942), bispectrum and gated (both 0.941). All equivariant methods cluster within $\sim$0.5pp of each other.

Scaling behavior. As parameters increase from the smallest to the largest configuration, standard, norm, and Fourier-ELU all degrade (by 1.8, 2.6, and 0.9 AUC points respectively), while the bispectrum improves (+0.9 points, reaching 0.943 at 920K params). This suggests that the bispectrum's complete invariant map scales more gracefully with capacity.

Data scaling (3 seeds). At 1% data the bispectrum leads all methods (0.912 ± 0.025), outperforming NormReLU (0.873), Fourier-ELU (0.876), and gate (0.884). By 10% all equivariant methods converge ($\sim$0.941–0.945). The bispectrum's data-efficiency advantage narrows as data increases, consistent with the hypothesis that complete invariants substitute for missing training examples in the low-data regime.

E OrganMNIST3D: Extended Results

E.1 Published baselines

Table 9 situates our results against published methods on the same benchmark. Our models are deliberately small ($\leq$463K parameters) to enable controlled ablation of the pooling mechanism. The relevant comparison is against equivariant methods at a similar scale.

Table 9: Published baselines on OrganMNIST3D for context. Accuracy differences with respect to large-scale models reflect the deliberate parameter constraint of our ablation study.

Method	Reference	Params	ACC	AUC
ResNet-18 + 3D	Yang et al. 2023	33M	0.907	0.996
ResNet-18 + ACS	Yang et al. 2023	11M	0.900	0.994
ILPOResNet-50	Zhemchuzhnikov and Grudinin [2024]	38K	0.879	0.992
EquiLoPO (local train.)	Zhemchuzhnikov and Grudinin [2025]	418K	0.866	0.991
SE3MovFrNet	Sangalli et al. [2023]	—	0.745	—
Regular SE(3) conv	Kuipers and Bekkers [2023]	172K	0.698	—

E.2 Rotation analysis

Figure 7 shows that all equivariant models achieve identical performance on the original and octahedrally rotated test sets ($\sigma_{\text{rot}} \approx 0$), while the augmentation-trained standard CNN degrades under rotation.

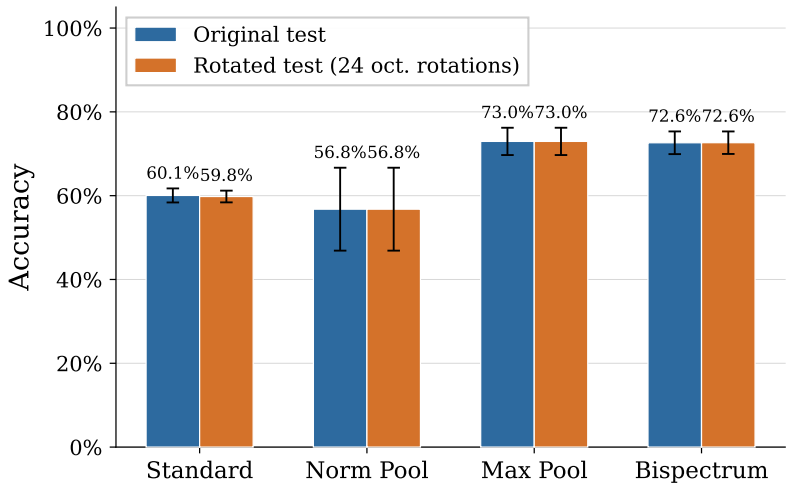

Figure 7: OrganMNIST3D test accuracy on the original and octahedrally rotated test sets. Error bars show standard deviation over three seeds. All equivariant models achieve identical performance under rotation ($\sigma_{\text{rot}} \approx 0$), while the augmentation-trained standard CNN degrades.

E.3 Wider channels

To test whether the advantage of bispectral completeness persists at higher capacity, we run max pool and bispectrum at channel widths $(8, 16)$ and $(16, 32)$ with full training data.

At $(8, 16)$ the two methods are equivalent ($74.5 \pm 0.6\%$ bispectrum vs. $74.3 \pm 1.5\%$ max pool, both at $\sim$1.5–1.7M parameters). At $(16, 32)$ max pool reaches $78.5 \pm 3.2\%$ while bispectrum drops to $68.5 \pm 3.9\%$—the 6.3M parameter bispectrum model likely overfits the 971 training samples. This suggests that the bispectrum's completeness advantage is most pronounced in the data-limited rather than the parameter-limited regime, consistent with the data efficiency findings in the main text.

Table 10: OrganMNIST3D accuracy at wider channel configurations (mean $\pm$ std over 3 seeds, full training data).

Channels	Params	Max Pool	Bispectrum
$(4, 8)$	$\sim$375K / 463K	$.730 \pm .033$	$.726 \pm .027$
$(8, 16)$	$\sim$1.5M / 1.7M	$.743 \pm .015$	$.745 \pm .006$
$(16, 32)$	$\sim$6.0M / 6.3M	$\mathbf{.785 \pm .032}$	$.685 \pm .039$

F Experimental Details

F.1 Hyperparameters

Table 11 summarises the training hyperparameters for all three experimental settings.

Table 11: Training hyperparameters across experiments.

	PCam	OrganMNIST3D	Sph. MNIST
Optimizer	AdamW	AdamW	AdamW
Learning rate	10^{-3}	10^{-3}	10^{-3}
Weight decay	10^{-4}	10^{-4}	10^{-4}
Batch size	64–256	16–64	256
Scheduler	Cosine warm restarts	—	—
Max epochs	50	100	50
Early stopping	Val AUC, pat. 10	Val AUC, pat. 15	Val ACC, pat. 10
Seeds	42, 123, 456†	42, 123, 456	42, 123, 456
Mixed precision	Yes	No	No

†Pareto figure uses seed 42; matched-parameter table and data efficiency use seeds 42, 123, 456.

F.2 Hardware

All experiments were run on a single NVIDIA H100 GPU (80 GB). Training wall-clock times range from $\sim$5 minutes (OrganMNIST3D, 971 samples) to $\sim$20 minutes (PCam, 26K samples at 10% data). Spherical MNIST bispectrum computation (forward pass through `SO3onS2` at $\ell_{\max} = 15$) takes $\sim$0.3s per batch of 64 on GPU.

F.3 Data preprocessing

PCam. PatchCamelyon [Veeling et al., 2018] consists of 327,680 RGB patches of size 96×96 pixels extracted from whole-slide images of lymph node sections [Ehteshami Bejnordi et al., 2017]. We use the standard train/val/test split. Images are loaded from HDF5 files and normalised to $[0, 1]$. The standard (non-equivariant) baseline uses random rotation and reflection augmentation during training; all equivariant models are trained without augmentation.

OrganMNIST3D. OrganMNIST3D [Yang et al., 2023] contains 1,742 abdominal CT organ volumes of size $28 \times 28 \times 28$ voxels with 11 organ classes (train/val/test = 971/161/610). Volumes are normalised to $[0, 1]$. The standard baseline uses random octahedral augmentation (one of 24 rotations per sample); equivariant models are trained without augmentation. Kernel rotations are implemented as exact voxel permutations on the 3^3 grid, avoiding interpolation artifacts.

Spherical MNIST. MNIST digits are projected onto the sphere via stereographic projection following Cohen et al. [2018] and sampled on a 64×128 equiangular grid. The selective bispectrum is computed at $\ell_{\max} = 15$, producing 384 complex invariant coefficients (768 real features after taking log-magnitude), which are fed to a 3-layer MLP: Linear(768, 256)–BN–ReLU–Linear(256, 128)–BN–ReLU–Linear(128, 10). The power spectrum MLP uses the same structure with hidden=128: Linear(16, 128)–BN–ReLU–Linear(128, 64)–BN–ReLU–Linear(64, 10). The standard CNN has 4 convolutional layers (16–32–64–64 channels) followed by adaptive average pooling and a Linear(1024, 128)–ReLU–Linear(128, 10) classifier.

F.4 Aggregation procedures

PCam. The Pareto figure (Figure 2) uses a single seed (42) across the full growth-rate sweep. The matched-parameter table (Table 3) and data efficiency analysis use 3 seeds (42, 123, 456) per configuration.

OrganMNIST3D. All results in Table 4 report mean $\pm$ std over 3 seeds (42, 123, 456). Rotation robustness (σ_{rot}) is the standard deviation of test accuracy across all 24 octahedral rotations of the test set. The data efficiency analysis (Figure 3 in the main text) uses 3 seeds per data fraction.

Spherical MNIST. All results in Table 5 report mean $\pm$ std over 3 seeds (42, 123, 456). Rotation robustness (σ_{rot}) is measured by applying 10 random $\mathrm{SO}(3)$ rotations to the entire test set and computing the standard deviation of accuracy, yielding 30 measurements per model (10 rotations $\times$ 3 seeds).

F.5 Dataset licenses

- **PCam**: CC0 1.0 Universal (public domain).
- **OrganMNIST3D**: CC BY 4.0.
- **Spherical MNIST**: derived from MNIST, which is publicly available for research use.

G API Example

```
import torch
from bispectrum import CnonCn, SO2onDisk, OctaonOcta

bsp = CnonCn(n=8, selective=True)
f = torch.randn(16, 8)
beta = bsp(f)
f_rec = bsp.invert(beta)
```

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and introduction list four contributions (selective bispectrum theory, completeness proof, `nn.Module` library, and experiments on three benchmarks) that are substantiated in Sections 3–4 and Appendices A–B.

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Section 6 contains a dedicated Limitations paragraph discussing the genericity assumption, bootstrap verification scope ($L \leq 100$), octahedral inversion recovery rate, and missing experimental comparisons. Appendix A includes a "Limitations and open questions" subsection.

   Guidelines:

   - The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Theorem 1 states all assumptions (genericity, gauge-fixing) and a proof sketch appears in Section 3.4. The complete proof with 10 lemmas and 6 propositions is in Appendix A. The octahedral inversion gap analysis is in Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All hyperparameters are in Appendix F (Table 11), architectures in Appendices D–E, optimizer and scheduler details in Table 11, and random seeds are specified. The code is provided as an anonymous repository.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

(d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: An anonymous code repository is provided with the submission. All three datasets (PCam, OrganMNIST3D, Spherical MNIST) are publicly available. Sweep scripts and training code are included.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Sections 4.1–4.3 describe experimental setups. Full hyperparameters are in Table 11 (Appendix F), architecture details in Appendices D–E, data preprocessing and splits in Appendix F.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Tables 4–5 (OrganMNIST3D, Spherical MNIST) report mean $\pm$ standard deviation over 3 seeds. Table 3 (PCam) uses a single seed for the Pareto sweep, which is noted in Appendix F. The aggregation procedure is described in Appendix F.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix F states that all experiments were run on a single NVIDIA H100 GPU (80 GB). Wall-clock times are available in the code repository.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

Answer: [Yes]

Justification: The work develops mathematical invariants and a software library evaluated on publicly available datasets. No human subjects, private data, or dual-use concerns are involved.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: This is foundational mathematical and software research on group-theoretic invariants. We do not foresee direct negative societal impacts.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The released asset is a deterministic mathematical library for computing group-theoretic invariants. It poses no risk for misuse.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All datasets are cited (PCam: CC0 1.0, OrganMNIST3D: CC BY 4.0, MNIST: public domain for research). Dataset licenses are listed in Appendix F.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The library is released under the MIT license with documentation, API reference, and usage examples. Section 3 describes the design and interface.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This work does not involve crowdsourcing or human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This work does not involve human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLMs are not part of the core methodology of this work.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.